

## Deep Learning Assignment 2

Name: Santosh Srinivas Ravichandran

Net ID: sr6411

1.

a)

$$L = 0.5 * (a * w_1^2 + b * w_2^2)$$

$$\frac{\partial L}{\partial w_1} = 0.5 * a * 2 * w_1 = a * w_1$$

$$\frac{\partial L}{\partial w_2} = 0.5 * b * 2 * w_2 = b * w_2$$

$$\nabla L(w) = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \frac{\partial L}{\partial w_2} \end{bmatrix}$$

$$\nabla L(w) = \begin{bmatrix} a * w_1 \\ b * w_2 \end{bmatrix}$$

To derive the weights where the function is minimum, we set the gradient to 0 and find the weights where this occurs.

$$\text{Setting } \frac{\partial L}{\partial w_1} = 0 \text{ and } \frac{\partial L}{\partial w_2} = 0$$

$$\frac{\partial L}{\partial w_1} = a(w_1^*) = 0$$

$$\Rightarrow w_1^* = 0$$

$$\frac{\partial L}{\partial w_2} = b(w_2^*) = 0$$

$$\Rightarrow w_2^* = 0$$

Thus the weights,  $w_1^*$  and  $w_2^*$  that achieve the minimum value of L are 0 and 0.

$w^* = (w_1^*, w_2^*) = (0, 0)$  are the weights that achieve the minimum value of L.

B )

Gradient Descent formula:

$$w_i(t + 1) = w_i(t) - \alpha * \frac{\partial L}{\partial w_i(t)} \text{ where } \alpha \text{ is the learning rate.}$$

Thus for  $w_1$ , that is for  $i=1$

$$w_1(t + 1) = w_1(t) - \alpha * \frac{\partial L}{\partial w_1(t)}$$

Substituting  $\frac{\partial L}{\partial w_1(t)} = a * w_1(t)$  in the above equation

$$w_1(t + 1) = w_1(t) - \alpha * a * w_1(t)$$

$$w_1(t + 1) = (1 - \alpha * a) * w_1(t)$$

$$w_1(t + 1) = \rho_1 * w_1(t) \text{ where } \rho_1 = (1 - \alpha * a)$$

Similarly for  $w_2$ ,

$$w_2(t + 1) = w_2(t) - \alpha * \frac{\partial L}{\partial w_2(t)}$$

Substituting  $\frac{\partial L}{\partial w_2(t)} = b * w_2(t)$  in the above equation

$$w_2(t + 1) = w_2(t) - \alpha * b * w_2(t)$$

$$w_2(t + 1) = (1 - \alpha * b) * w_2(t)$$

$$w_2(t + 1) = \rho_2 * w_2(t) \text{ where } \rho_2 = (1 - \alpha * b)$$

C )

$$\text{Since } L = 0.5 * (a * w_1^2 + b * w_2^2)$$

This is a convex shaped function like a bowl, with its minima at (0,0)

$$w_1 \text{ updations is like this: } w_1(t + 1) = \rho_1 * w_1(t) \text{ where } \rho_1 = (1 - \alpha * a)$$

$$w_2 \text{ updations is like this: } w_2(t + 1) = \rho_2 * w_2(t) \text{ where } \rho_2 = (1 - \alpha * b)$$

Thus if  $0 < \rho_1 < 1$  and  $0 < \rho_2 < 1$ , each update for both  $w_1$  and  $w_2$  will bring us closer to the minima directly. This is because multiplying  $w_1(t)$ ,  $w_2(t)$  by a number less than 1 ( $\rho_1$ ,  $\rho_2$ ) will iteratively take it zero which is the minima.

$$\Rightarrow 0 < 1 - \alpha * a < 1 \text{ and } 0 < 1 - \alpha * b < 1 \text{ -----equations 1}$$

$$\Rightarrow 0 < \alpha * a < 1 \text{ and } 0 < \alpha * b < 1$$

$\Rightarrow 0 < \alpha < 1/a \text{ and } 0 < \alpha < 1/b$  ( assuming a,b are non-negative )

Whichever of  $1/a$  and  $1/b$  is lesser is chosen as the upper bound for the learning rate while 0 is the lower bound.

$\Rightarrow 0 < \alpha < \min(1/a, 1/b)$

Thus when a and b are non negative,  $0 < \alpha < \min(1/a, 1/b)$  does the job of convergence.

If a is negative,  $\Rightarrow 1/a < \alpha < 0$  from equation 1 ( reversing inequality due to negative sign)

This means mathematically learning rate of negative value can do the job of directly converging. But in machine learning negative learning rates are not the most intuitive. Thus when a or b is negative,  $0 < \alpha < \min(1/a, 1/b)$  doesn't hold as the learning rate becomes negative.

However, learning rates of positive value such as between 0.0001 and 1 can lead to convergence. The convergence may not be direct as shown in the case where a and b are positive while we set learning rate to  $0 < \alpha < \min(1/a, 1/b)$  but with several iterations it is shown to work.

In addition, learning rates can be changed with time like in adaptive scheduling, where

$\alpha = k/(t)^{0.5}$  where t is the time

Thus a learning rate value initialized with 1 can become 0.0001 with several iterations.

As shown in the demo for quadratic loss functions:

Newton's Method can have learning rate of 0.5 to lead to convergence.

For momentum, a learning rate of 1.2 can lead to convergence effectively.

For Adam, a learning rate of 1.6 can lead to convergence.

Thus, to summarize, a learning rate of 0.0001(approx) to 1.5( approx) can lead to convergence for our loss function depending on the optimisation algorithm. A very large value such as 100 etc will definitely not lead to convergence as the weights overshoot the optimum continuously.

D)

A scenario where a slow convergence can occur is when the a/b ratio is very large.

Since  $\frac{\partial L}{\partial w_1} = a(w_1^*)$  and  $\frac{\partial L}{\partial w_2} = b(w_2^*)$  The imbalance in the values of a and b means that the optimization algorithm is highly sensitive to changes in w1 but not as sensitive to changes in w2.

Thus this can lead to oscillations in the optimization process, where the algorithm continually overshoots and undershoots the optimal w1 value, leading to a slow and oscillatory convergence.

2.

a)

Sobel Filters:

G\_X =

-1	0	1
-2	0	2
-1	0	1

AND

G\_Y =

1	2	1
0	0	0
-1	-2	-1

The filter with weights G\_X above is used for vertical edge detection and outputs a high value when the image has a vertical edge. This is because to the left and right side of the vertical edge, there is a gradient in pixel values as the edge separates the foreground pixels from the background pixels. This is detected with the weights G\_X as the right most column is [1,2,1] and

the leftmost column is  $[-1, -2, -1]$ . When there is no edge in the image, that is when there is uniform pixel intensity, the filter outputs a low value.

The filter with weights  $G_Y$  above is used for horizontal edge detection and outputs a high value when the image has a horizontal edge. This is because to the top and bottom side of the horizontal edge, there is a gradient in pixel values as the edge separates the foreground pixels from the background pixels. This is detected with the weights  $G_Y$  as the top most row is  $[1, 2, 1]$  and the bottom most row is  $[-1, -2, -1]$ . When there is no edge in the image, that is when there is uniform pixel intensity, the filter outputs a low value.

b)

Gaussian Filter:

$$W = \begin{array}{|c|c|c|} \hline 1/16 & 1/8 & 1/16 \\ \hline 1/8 & 1/4 & 1/8 \\ \hline 1/16 & 1/8 & 1/16 \\ \hline \end{array}$$

To create a blurring filter for a 2D image, the weights must emphasize on the smoothing of pixel values across the image. The Gaussian filter provides a weighted average of the pixel values within the filter region with the property that they give higher weights to pixels closer to the center and lower weights to pixels farther away, which effectively smooths out the finer details of the image making it appear blurred.

C )

Sobel Filter:

$$W = \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

This filter can work for horizontal sharpening because after convolving with the filter, the horizontal edges and features in the images are emphasized. This is because the pixel values to the top and bottom of a horizontal edge differ ( foreground and background separation ) and this filter captures this difference due to positive weights on the topmost row and negative

weights on the bottom most row. Thus, this design aims to enhance the contrast between pixel values along horizontal lines or edges resulting in horizontal sharpening.

D)

Gaussian Filter:

W =

1/16	1/8	1/16
1/8	1/4	1/8
1/16	1/8	1/16

The above Gaussian filter has the property of providing weighted averaging of pixel values within a local neighborhood around each pixel. The Gaussian distribution ensures that nearby pixels receive higher weights, while more distant pixels receive lower weights. This weighted averaging helps to reduce the impact of noisy pixel values. Noise, which is often characterized by high-frequency variations in pixel values, gets smoothed out as the filter emphasizes the nearby, more important pixel values.

3.A)

The IoU (Intersection over Union) metric between two bounding boxes can be defined as the ratio of the area of their intersection to the area of their union.

The area of intersection and the area of union, both are non negative.

Thus,  $\text{IoU} = (\text{area of intersection} / \text{area of union}) \geq 0$

The area of intersection is always less or equal to the area of Union.

Thus,  $\text{IoU} = (\text{area of intersection} / \text{area of union}) \leq 1$

When two bounding boxes don't overlap at all, IoU is 0 as area of intersection is 0

When two bounding boxes fully overlap, IoU is 1 as area of intersection = area of union

When two bounding boxes partially overlap, IoU is greater than 0 but less than 1 as area of intersection < area of union

Thus IoU will be a real number in the range  $[0, 1]$

b)

Consider the IoU metric expressed as a function of the coordinates of the bounding boxes.

When the top-left corner of one of the bounding boxes changes, the top-left corner of the intersection bounding box may or may not change, depending on the new position of the bounding box. This introduces a non-differentiable, piecewise behavior.

If the top-left corner of one bounding box moves in a way that it doesn't affect the overlap with the other box, the IoU remains the same. However, if the movement causes the boxes to overlap, the IoU starts to increase from 0 to a small value and then continues to increase to 1 as the overlap becomes larger. After reaching 1, if the movement continues, the IoU decreases again, eventually reaching 0 as the two bounding boxes no longer overlap. This piecewise behavior is characteristic of non-differentiability. This non-differentiability arises due to the sudden changes in the IoU value as the relative positions of the bounding boxes change.

Similar arguments can be applied to the other three corners of the bounding boxes, and the non-differentiability issue persists across all corners and edges.

Because of its non-differentiability, it's challenging to use the IoU directly as a loss function for training neural networks. This is because gradient-based optimization methods, such as backpropagation, require the loss function to be differentiable, as gradients are used to update model parameters during training.



# AlexNet

In this problem, you are asked to train a deep convolutional neural network to perform image classification. In fact, this is a slight variation of a network called *AlexNet*. This is a landmark model in deep learning, and arguably kickstarted the current (and ongoing, and massive) wave of innovation in modern AI when its results were first presented in 2012. AlexNet was the first real-world demonstration of a *deep* classifier that was trained end-to-end on data and that outperformed all other ML models thus far.

We will train AlexNet using the [CIFAR10](#) dataset, which consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. The classes are: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck.

A lot of the code you will need is already provided in this notebook; all you need to do is to fill in the missing pieces, and interpret your results.

**Warning :** AlexNet takes a good amount of time to train (~1 minute per epoch on Google Colab). So please budget enough time to do this homework.

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.optim.lr_scheduler import _LRScheduler
import torch.utils.data as data

import torchvision.transforms as transforms
import torchvision.datasets as datasets

from sklearn import decomposition
from sklearn import manifold
```

```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np

import copy
import random
import time

SEED = 1234

random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

```

## Loading and Preparing the Data

Our dataset is made up of color images but three color channels (red, green and blue), compared to MNIST's black and white images with a single color channel. To normalize our data we need to calculate the means and standard deviations for each of the color channels independently, and normalize them.

```

ROOT = '.data'
train_data = datasets.CIFAR10(root = ROOT,
                              train = True,
                              download = True)

Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to
.data/cifar-10-python.tar.gz
100%|██████████| 170498071/170498071 [00:10<00:00, 16095863.59it/s]
Extracting .data/cifar-10-python.tar.gz to .data

# Compute means and standard deviations along the R,G,B channel

means = train_data.data.mean(axis = (0,1,2)) / 255
stds = train_data.data.std(axis = (0,1,2)) / 255

```

Next, we will do data augmentation. For each training image we will randomly rotate it (by up to 5 degrees), flip/mirror with probability 0.5, shift by +/-1 pixel. Finally we will normalize each color channel using the means/stds we calculated above.

```

train_transforms = transforms.Compose([
    transforms.RandomRotation(5),
    transforms.RandomHorizontalFlip(0.5),

```

```

        transforms.RandomCrop(32, padding = 2),
        transforms.ToTensor(),
        transforms.Normalize(mean = means,
                               std = stds)
    ])

test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean = means,
                          std = stds)
])

```

Next, we'll load the dataset along with the transforms defined above.

We will also create a validation set with 10% of the training samples. The validation set will be used to monitor loss along different epochs, and we will pick the model along the optimization path that performed the best, and report final test accuracy numbers using this model.

```

train_data = datasets.CIFAR10(ROOT,
                               train = True,
                               download = True,
                               transform = train_transforms)

test_data = datasets.CIFAR10(ROOT,
                              train = False,
                              download = True,
                              transform = test_transforms)

Files already downloaded and verified
Files already downloaded and verified

VALID_RATIO = 0.9

n_train_examples = int(len(train_data) * VALID_RATIO)
n_valid_examples = len(train_data) - n_train_examples

train_data, valid_data = data.random_split(train_data,
                                           [n_train_examples,
                                           n_valid_examples])

valid_data = copy.deepcopy(valid_data)
valid_data.dataset.transform = test_transforms

```

Now, we'll create a function to plot some of the images in our dataset to see what they actually look like.

Note that by default PyTorch handles images that are arranged [channel, height, width], but matplotlib expects images to be [height, width, channel], hence we need to permute the dimensions of our images before plotting them.

```

def plot_images(images, labels, classes, normalize = False):
    n_images = len(images)

    rows = int(np.sqrt(n_images))
    cols = int(np.sqrt(n_images))

    fig = plt.figure(figsize = (10, 10))

    for i in range(rows*cols):
        ax = fig.add_subplot(rows, cols, i+1)

        image = images[i]

        if normalize:
            image_min = image.min()
            image_max = image.max()
            image.clamp_(min = image_min, max = image_max)
            image.add_(-image_min).div_(image_max - image_min + 1e-5)

        ax.imshow(image.permute(1, 2, 0).cpu().numpy())
        ax.set_title(classes[labels[i]])
        ax.axis('off')

```

One point here: `matplotlib` is expecting the values of every pixel to be between  $[0, 1]$ , however our normalization will cause them to be outside this range. By default `matplotlib` will then clip these values into the  $[0, 1]$  range. This clipping causes all of the images to look a bit weird - all of the colors are oversaturated. The solution is to normalize each image between  $[0, 1]$ .

```

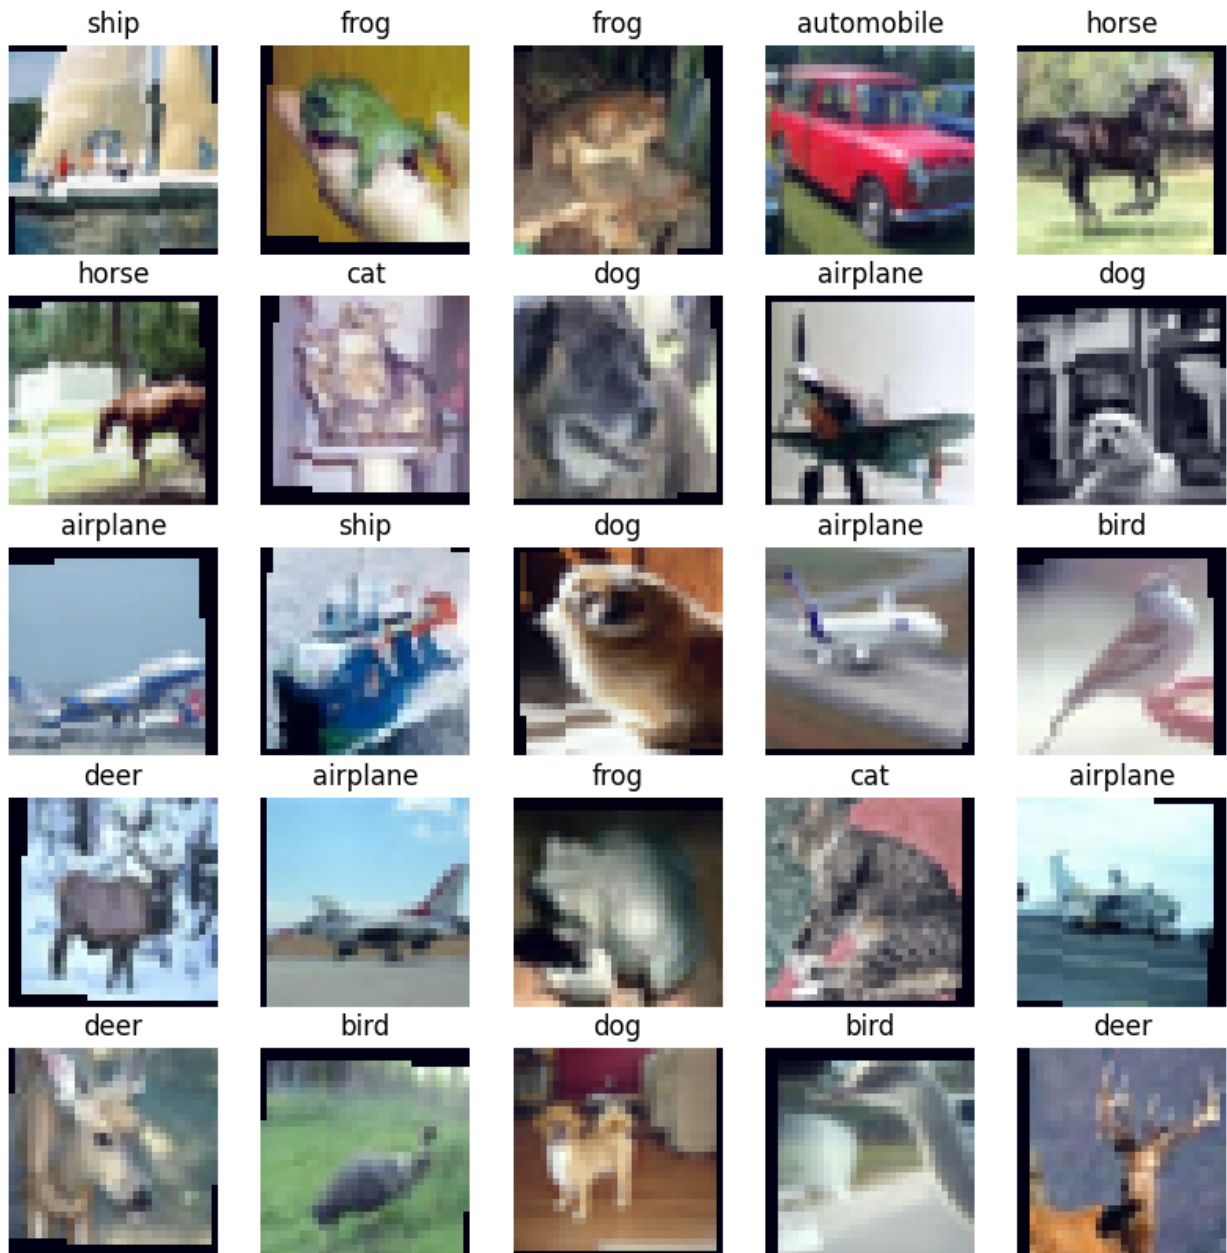
N_IMAGES = 25

images, labels = zip(*[(image, label) for image, label in
                        [train_data[i] for i in range(N_IMAGES)]])

classes = test_data.classes

plot_images(images, labels, classes, normalize = True)

```



We'll be normalizing our images by default from now on, so we'll write a function that does it for us which we can use whenever we need to renormalize an image.

```
def normalize_image(image):
    image_min = image.min()
    image_max = image.max()
    image.clamp_(min = image_min, max = image_max)
    image.add_(-image_min).div_(image_max - image_min + 1e-5)
    return image
```

The final bit of the data processing is creating the iterators. We will use a large. Generally, a larger batch size means that our model trains faster but is a bit more susceptible to overfitting.

```

# Q1: Create data loaders for train_data, valid_data, test_data
# Use batch size 256

#import utils

BATCH_SIZE = 256

train_iterator = torch.utils.data.DataLoader(
    train_data,
    batch_size=BATCH_SIZE,
    shuffle=True,
    num_workers=4,
)

valid_iterator = torch.utils.data.DataLoader(
    valid_data,
    batch_size=BATCH_SIZE,
    shuffle=False,
    num_workers=4,
)

test_iterator = torch.utils.data.DataLoader(
    test_data,
    batch_size=BATCH_SIZE,
    shuffle=False,
    num_workers=4,
)

```

## Defining the Model

Next up is defining the model.

AlexNet will have the following architecture:

- There are 5 2D convolutional layers (which serve as *feature extractors*), followed by 3 linear layers (which serve as the *classifier*).
- All layers (except the last one) have ReLU activations. (Use `inplace=True` while defining your ReLUs.)
- All convolutional filter sizes have kernel size 3 x 3 and padding 1.
- Convolutional layer 1 has stride 2. All others have the default stride (1).
- Convolutional layers 1, 2, and 5 are followed by a 2D maxpool of size 2.
- Linear layers 1 and 2 are preceded by Dropouts with Bernoulli parameter 0.5.

- For the convolutional layers, the number of channels is set as follows. We start with 3 channels and then proceed like this:

–  $3 \rightarrow 64 \rightarrow 192 \rightarrow 384 \rightarrow 256 \rightarrow 256$

In the end, if everything is correct you should get a feature map of size  $2 \times 2 \times 256 = 1024$ .

- For the linear layers, the feature sizes are as follows:

–  $1024 \rightarrow 4096 \rightarrow 4096 \rightarrow 10$ .

(The 10, of course, is because 10 is the number of classes in CIFAR-10).

```
class AlexNet(nn.Module):
    def __init__(self, output_dim):
        super().__init__()

        self.features = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=3, stride=2, padding=1),
            nn.MaxPool2d(kernel_size=2),
            nn.ReLU(inplace=True),

            nn.Conv2d(64, 192, kernel_size=3, padding=1),
            nn.MaxPool2d(kernel_size=2),
            nn.ReLU(inplace=True),

            nn.Conv2d(192, 384, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            nn.Conv2d(384, 256, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),

            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.MaxPool2d(kernel_size=2),
            nn.ReLU(inplace=True),
        )

        # Classifier (Linear Layers)
        self.classifier = nn.Sequential(
            nn.Dropout(0.5),
            nn.Linear(256 * 2 * 2, 4096),
            nn.ReLU(inplace=True),
            nn.Dropout(0.5),
            nn.Linear(4096, 4096),
            nn.ReLU(inplace=True),
            nn.Linear(4096, output_dim)
        )
```

```
def forward(self, x):
    x = self.features(x)
    h = x.view(x.shape[0], -1)
    x = self.classifier(h)
    return x, h
```

We'll create an instance of our model with the desired amount of classes.

```
OUTPUT_DIM = 10
model = AlexNet(OUTPUT_DIM)
```

## Training the Model

We first initialize parameters in PyTorch by creating a function that takes in a PyTorch module, checking what type of module it is, and then using the `nn.init` methods to actually initialize the parameters.

For convolutional layers we will initialize using the *Kaiming Normal* scheme, also known as *He Normal*. For the linear layers we initialize using the *Xavier Normal* scheme, also known as *Glorot Normal*. For both types of layer we initialize the bias terms to zeros.

```
def initialize_parameters(m):
    if isinstance(m, nn.Conv2d):
        nn.init.kaiming_normal_(m.weight.data, nonlinearity = 'relu')
        nn.init.constant_(m.bias.data, 0)
    elif isinstance(m, nn.Linear):
        nn.init.xavier_normal_(m.weight.data, gain =
nn.init.calculate_gain('relu'))
        nn.init.constant_(m.bias.data, 0)
```

We apply the initialization by using the model's `apply` method. If your definitions above are correct you should get the printed output as below.

```
model.apply(initialize_parameters)

AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    (1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (2): ReLU(inplace=True)
    (3): Conv2d(64, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): ReLU(inplace=True)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1),
```



```
padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (11): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (12): ReLU(inplace=True)
)
(classifier): Sequential(
  (0): Dropout(p=0.5, inplace=False)
  (1): Linear(in_features=1024, out_features=4096, bias=True)
  (2): ReLU(inplace=True)
  (3): Dropout(p=0.5, inplace=False)
  (4): Linear(in_features=4096, out_features=4096, bias=True)
  (5): ReLU(inplace=True)
  (6): Linear(in_features=4096, out_features=10, bias=True)
)
)
```

We then define the loss function we want to use, the device we'll use and place our model and criterion on to our device.

```
optimizer = optim.Adam(model.parameters(), lr = 1e-3)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
criterion = nn.CrossEntropyLoss()

model = model.to(device)
criterion = criterion.to(device)

# This is formatted as code
```

We define a function to calculate accuracy...

```
def calculate_accuracy(y_pred, y):
    top_pred = y_pred.argmax(1, keepdim = True)
    correct = top_pred.eq(y.view_as(top_pred)).sum()
    acc = correct.float() / y.shape[0]
    return acc
```

As we are using dropout we need to make sure to "turn it on" when training by using `model.train()`.

```
def train(model, iterator, optimizer, criterion, device):
    epoch_loss = 0
    epoch_acc = 0
```

```

model.train()

for (x, y) in iterator:
    x = x.to(device)
    y = y.to(device)

    optimizer.zero_grad()

    y_pred, _ = model(x)

    loss = criterion(y_pred, y)

    acc = calculate_accuracy(y_pred, y)

    loss.backward()

    optimizer.step()

    epoch_loss += loss.item()
    epoch_acc += acc.item()

return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

We also define an evaluation loop, making sure to "turn off" dropout with `model.eval()`.

```

def evaluate(model, iterator, criterion, device):
    epoch_loss = 0
    epoch_acc = 0

    model.eval()

    with torch.no_grad():
        for (x, y) in iterator:
            x = x.to(device)
            y = y.to(device)

            y_pred, _ = model(x)

            loss = criterion(y_pred, y)

            acc = calculate_accuracy(y_pred, y)

            epoch_loss += loss.item()
            epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

Next, we define a function to tell us how long an epoch takes.

```
def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)
    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs
```

Then, finally, we train our model.

Train it for 25 epochs (using the train dataset). At the end of each epoch, compute the validation loss and keep track of the best model. You might find the command `torch.save` helpful.

At the end you should expect to see validation losses of ~76% accuracy.

```
# Q3: train your model here for 25 epochs.
# Print out training and validation loss/accuracy of the model after
each epoch
# Keep track of the model that achieved best validation loss thus far.

EPOCHS = 25

# Fill training code here

best_valid_loss = float('inf') # Initialize with a large value
best_model = None

for epoch in range(EPOCHS):
    start_time = time.time()

    # Training
    train_loss, train_acc = train(model, train_iterator, optimizer,
    criterion, device)

    # Validation
    valid_loss, valid_acc = evaluate(model, valid_iterator, criterion,
    device)

    end_time = time.time()

    # Calculate the time for the current epoch
    epoch_mins, epoch_secs = epoch_time(start_time, end_time)

    # Print the results for the current epoch
    print(f'Epoch: {epoch + 1:02}')
    print(f'\tTime: {epoch_mins}m {epoch_secs}s')
    print(f'\tTrain Loss: {train_loss:.3f} | Train Acc:
    {train_acc*100:.2f}%')
    print(f'\tValid Loss: {valid_loss:.3f} | Valid Acc:
    {valid_acc*100:.2f}%')
```

```
# Check if this model has the best validation loss
if valid_loss < best_valid_loss:
    best_valid_loss = valid_loss
    best_model = model.state_dict() # Save the model's state_dict
```

*# After training, you can use 'best\_model' to load the best model.*

```
Epoch: 01
  Time: 0m 15s
  Train Loss: 2.349 | Train Acc: 22.82%
  Valid Loss: 1.567 | Valid Acc: 42.05%
Epoch: 02
  Time: 0m 7s
  Train Loss: 1.519 | Train Acc: 43.57%
  Valid Loss: 1.342 | Valid Acc: 51.10%
Epoch: 03
  Time: 0m 7s
  Train Loss: 1.358 | Train Acc: 50.59%
  Valid Loss: 1.258 | Valid Acc: 55.01%
Epoch: 04
  Time: 0m 8s
  Train Loss: 1.256 | Train Acc: 54.89%
  Valid Loss: 1.159 | Valid Acc: 59.07%
Epoch: 05
  Time: 0m 7s
  Train Loss: 1.187 | Train Acc: 57.47%
  Valid Loss: 1.117 | Valid Acc: 59.93%
Epoch: 06
  Time: 0m 7s
  Train Loss: 1.111 | Train Acc: 60.45%
  Valid Loss: 1.044 | Valid Acc: 63.01%
Epoch: 07
  Time: 0m 7s
  Train Loss: 1.056 | Train Acc: 63.03%
  Valid Loss: 0.997 | Valid Acc: 65.51%
Epoch: 08
  Time: 0m 7s
  Train Loss: 1.015 | Train Acc: 64.42%
  Valid Loss: 0.952 | Valid Acc: 66.54%
Epoch: 09
  Time: 0m 7s
  Train Loss: 0.969 | Train Acc: 66.04%
  Valid Loss: 0.902 | Valid Acc: 68.51%
Epoch: 10
  Time: 0m 7s
  Train Loss: 0.928 | Train Acc: 67.56%
  Valid Loss: 0.870 | Valid Acc: 70.22%
Epoch: 11
  Time: 0m 7s
  Train Loss: 0.903 | Train Acc: 68.43%
```

```
Valid Loss: 0.869 | Valid Acc: 70.53%
Epoch: 12
Time: 0m 7s
Train Loss: 0.865 | Train Acc: 69.78%
Valid Loss: 0.817 | Valid Acc: 72.11%
Epoch: 13
Time: 0m 7s
Train Loss: 0.847 | Train Acc: 70.67%
Valid Loss: 0.860 | Valid Acc: 70.88%
Epoch: 14
Time: 0m 7s
Train Loss: 0.815 | Train Acc: 71.98%
Valid Loss: 0.808 | Valid Acc: 72.67%
Epoch: 15
Time: 0m 7s
Train Loss: 0.793 | Train Acc: 72.47%
Valid Loss: 0.780 | Valid Acc: 74.08%
Epoch: 16
Time: 0m 8s
Train Loss: 0.779 | Train Acc: 73.00%
Valid Loss: 0.803 | Valid Acc: 73.38%
Epoch: 17
Time: 0m 7s
Train Loss: 0.750 | Train Acc: 74.50%
Valid Loss: 0.775 | Valid Acc: 74.60%
Epoch: 18
Time: 0m 7s
Train Loss: 0.738 | Train Acc: 74.43%
Valid Loss: 0.772 | Valid Acc: 74.15%
Epoch: 19
Time: 0m 7s
Train Loss: 0.716 | Train Acc: 75.32%
Valid Loss: 0.758 | Valid Acc: 74.40%
Epoch: 20
Time: 0m 7s
Train Loss: 0.703 | Train Acc: 75.62%
Valid Loss: 0.743 | Valid Acc: 75.58%
Epoch: 21
Time: 0m 7s
Train Loss: 0.682 | Train Acc: 76.59%
Valid Loss: 0.737 | Valid Acc: 75.98%
Epoch: 22
Time: 0m 8s
Train Loss: 0.668 | Train Acc: 77.08%
Valid Loss: 0.746 | Valid Acc: 75.04%
Epoch: 23
Time: 0m 7s
Train Loss: 0.670 | Train Acc: 76.83%
Valid Loss: 0.744 | Valid Acc: 75.70%
```

```
Epoch: 24
  Time: 0m 7s
  Train Loss: 0.652 | Train Acc: 77.64%
  Valid Loss: 0.747 | Valid Acc: 75.94%
Epoch: 25
  Time: 0m 7s
  Train Loss: 0.631 | Train Acc: 78.37%
  Valid Loss: 0.711 | Valid Acc: 76.75%
```

## Evaluating the model

We then load the parameters of our model that achieved the best validation loss. You should expect to see ~75% accuracy of this model on the test dataset.

Finally, plot the confusion matrix of this model and comment on any interesting patterns you can observe there. For example, which two classes are confused the most?

```
# Q4: Load the best performing model, evaluate it on the test dataset,
and print test accuracy.
```

```
# Also, print out the confusion matrox.
```

```
# Load the best performing model,
best_model_state_dict = best_model
model.load_state_dict(best_model_state_dict)
```

```
<All keys matched successfully>
```

```
def get_predictions(model, iterator, device):
```

```
    model.eval()
```

```
    labels = []
    probs = []
```

```
    epoch_loss = 0
    epoch_acc = 0
```

```
# Q4: Fill code here.
```

```
with torch.no_grad():
    for (x, y) in iterator:
        x = x.to(device)
        y = y.to(device)
        y_pred, _ = model(x)
```

```
# Get predicted labels and their probabilities
```

```
_, predicted_labels = torch.max(y_pred, 1)
predicted_probs = torch.softmax(y_pred, dim=1)
```

```

        # Append true labels and predicted probabilities to their
        respective lists
        labels.append(y)
        probs.append(predicted_probs)

        # Calculate loss and accuracy for the current batch
        loss = criterion(y_pred, y)
        acc = calculate_accuracy(y_pred, y)

        # Accumulate batch loss and accuracy to compute epoch-
        level metrics
        epoch_loss += loss.item()
        epoch_acc += acc.item()

        # Calculate the average test loss and accuracy for the entire
        dataset
        test_loss = epoch_loss / len(iterator)
        test_acc = epoch_acc / len(iterator)
        print(f'Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}
        %')

        # Concatenate predicted probabilities and true labels for the
        entire dataset
        probs = torch.cat(probs, dim=0)
        labels = torch.cat(labels, dim=0)


        return labels, probs

labels, probs = get_predictions(model, test_iterator, device)
Test Loss: 0.688 | Test Acc: 76.51%
pred_labels = torch.argmax(probs, 1)

def plot_confusion_matrix(labels, pred_labels, classes):
    fig = plt.figure(figsize = (10, 10));
    ax = fig.add_subplot(1, 1, 1);
    cm = confusion_matrix(labels, pred_labels);
    cm = ConfusionMatrixDisplay(cm, display_labels = classes);
    cm.plot(values_format = 'd', cmap = 'Blues', ax = ax)
    plt.xticks(rotation = 20)

labels = labels.to('cpu')
pred_labels = pred_labels.to('cpu')
plot_confusion_matrix(labels, pred_labels, classes)

```



### Confusion Matrix comments:

See which diagonal element in the matrix ( same true and predicted label) has the least value, that must be the most confused class.

Thus, the two most confused classes as per the confusion matrix is cat and bird.

The third and fourth most confused classes are dog and deer.

The least confused classes is automobile.

### Conclusion

That's it! As a side project (this is not for credit and won't be graded), feel free to play around with different design choices that you made while building this network.

- Whether or not to normalize the color channels in the input.



- The learning rate parameter in Adam.
- The batch size.
- The number of training epochs.
- (and if you are feeling brave -- the AlexNet architecture itself.)

```
#Mounting google drive
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

## Defining the dataset

```
import os
import torch

from torchvision.io import read_image
from torchvision.ops.bboxes import masks_to_boxes
from torchvision import tv_tensors
from torchvision.transforms.v2 import functional as F

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root,
"PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root,
"PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
        img_path = os.path.join(self.root, "PNGImages",
self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks",
self.masks[idx])
        img = read_image(img_path)
        mask = read_image(mask_path)
        # instances are encoded as different colors
        obj_ids = torch.unique(mask)
        # first id is the background, so remove it
        obj_ids = obj_ids[1:]
        num_objs = len(obj_ids)

        # split the color-encoded mask into a set
        # of binary masks
        masks = (mask == obj_ids[:, None, None]).to(dtype=torch.uint8)

        # get bounding box coordinates for each mask
        boxes = masks_to_boxes(masks)

        # there is only one class
```

```

        labels = torch.ones((num_objs,), dtype=torch.int64)

        image_id = idx
        area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:,
0]))

        # suppose all instances are not crowd
        iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

        # Wrap sample and targets into torchvision tv_tensors:
        img = tv_tensors.Image(img)

        target = {}
        target["boxes"] = tv_tensors.BoundingBoxes(boxes,
format="XYXY", canvas_size=F.get_size(img))
        target["masks"] = tv_tensors.Mask(masks)
        target["labels"] = labels
        target["image_id"] = image_id
        target["area"] = area
        target["iscrowd"] = iscrowd

        if self.transforms is not None:
            img, target = self.transforms(img, target)

        return img, target

    def __len__(self):
        return len(self.imgs)

```

### Finetuning from a pretrained model (Option 1)

```

import torchvision
from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor

def get_fine_tuned_model_instance_segmentation(num_classes):
    # load an instance segmentation model pre-trained on COCO
    model =
torchvision.models.detection.maskrcnn_resnet50_fpn(weights="DEFAULT")

    # get number of input features for the classifier
    in_features = model.roi_heads.box_predictor.cls_score.in_features
    # replace the pre-trained head with a new one
    model.roi_heads.box_predictor = FastRCNNPredictor(in_features,
num_classes)

    # now get the number of input features for the mask classifier
    in_features_mask =
model.roi_heads.mask_predictor.conv5_mask.in_channels
    hidden_layer = 256

```

```

# and replace the mask predictor with a new one
model.roi_heads.mask_predictor = MaskRCNNPredictor(
    in_features_mask,
    hidden_layer,
    num_classes,
)

return model

```

Downloading some utility scripts and setting up data augmentation/transformations

```

# Download utility scripts for object detection from the PyTorch
Vision GitHub repository

os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detec
tion/engine.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detec
tion/utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detec
tion/coco_utils.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detec
tion/coco_eval.py")
os.system("wget
https://raw.githubusercontent.com/pytorch/vision/main/references/detec
tion/transforms.py")

# Since v0.15.0 torchvision provides `new Transforms API
<https://pytorch.org/vision/stable/transforms.html>`
# to easily write data augmentation pipelines for Object Detection and
Segmentation tasks.
#

# Let's write some helper functions for data augmentation /
# transformation:

from torchvision.transforms import v2 as T
import utils

def get_transform(train):
    transforms = []
    if train:
        transforms.append(T.RandomHorizontalFlip(0.5))
    transforms.append(T.ToDtype(torch.float, scale=True))
    transforms.append(T.ToPureTensor())
    return T.Compose(transforms)

```

### ***Training and validation for our Fine-Tuned Model***

```
from engine import train_one_epoch, evaluate

# train on the GPU or on the CPU, if a GPU is not available
device = torch.device('cuda') if torch.cuda.is_available() else
torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2
# use our dataset and defined transformations
dataset = PennFudanDataset('/content/drive/MyDrive/PennFudanPed',
get_transform(train=True))
dataset_test = PennFudanDataset('/content/drive/MyDrive/PennFudanPed',
get_transform(train=False))

# split the dataset in train and test set
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=4,
    collate_fn=utils.collate_fn
)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test,
    batch_size=1,
    shuffle=False,
    num_workers=4,
    collate_fn=utils.collate_fn
)

# get the model using our helper function
model = get_fine_tuned_model_instance_segmentation(num_classes)

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
```

```

        weight_decay=0.0005
    )

    # and a learning rate scheduler
    lr_scheduler = torch.optim.lr_scheduler.StepLR(
        optimizer,
        step_size=3,
        gamma=0.1
    )

    # let's train it for 10 epochs
    num_epochs = 10

    for epoch in range(num_epochs):
        # train for one epoch, printing every 10 iterations
        train_one_epoch(model, optimizer, data_loader, device, epoch,
            print_freq=10)
        # update the learning rate
        lr_scheduler.step()
        # evaluate on the test dataset
        evaluate(model, data_loader_test, device=device)

    print("That's it!")

```

```

/usr/local/lib/python3.10/dist-packages/torch/utils/data/
dataloader.py:557: UserWarning: This DataLoader will create 4 worker
processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to
create. Please be aware that excessive worker creation might get
DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.

```

```

    warnings.warn(_create_warning_msg(

```

```

Downloading:

```

```

"https://download.pytorch.org/models/maskrcnn_resnet50_fpn_coco-
bf2d0cle.pth" to

```

```

/root/.cache/torch/hub/checkpoints/maskrcnn_resnet50_fpn_coco-
bf2d0cle.pth

```

```

100%|██████████| 170M/170M [00:01<00:00, 92.1MB/s]

```

```

Epoch: [0] [ 0/60] eta: 0:12:14 lr: 0.000090 loss: 9.9941 (9.9941)
loss_classifier: 0.9717 (0.9717) loss_box_reg: 0.2146 (0.2146)
loss_mask: 8.7743 (8.7743) loss_objectness: 0.0326 (0.0326)
loss_rpn_box_reg: 0.0009 (0.0009) time: 12.2345 data: 3.3264 max
mem: 2151

```

```

Epoch: [0] [10/60] eta: 0:01:20 lr: 0.000936 loss: 2.5403 (4.3461)
loss_classifier: 0.5083 (0.5573) loss_box_reg: 0.2605 (0.2714)
loss_mask: 1.6730 (3.4892) loss_objectness: 0.0193 (0.0230)
loss_rpn_box_reg: 0.0039 (0.0052) time: 1.6123 data: 0.3083 max
mem: 3407

```

```

Epoch: [0] [20/60] eta: 0:00:43 lr: 0.001783 loss: 1.0655 (2.6353)

```

```
loss_classifier: 0.1359 (0.3474) loss_box_reg: 0.1742 (0.2197)
loss_mask: 0.5585 (2.0381) loss_objectness: 0.0197 (0.0249)
loss_rpn_box_reg: 0.0038 (0.0051) time: 0.5348 data: 0.0090 max
mem: 3407
Epoch: [0] [30/60] eta: 0:00:27 lr: 0.002629 loss: 0.6801 (1.9754)
loss_classifier: 0.0949 (0.2621) loss_box_reg: 0.1624 (0.2047)
loss_mask: 0.3543 (1.4814) loss_objectness: 0.0181 (0.0214)
loss_rpn_box_reg: 0.0042 (0.0059) time: 0.5279 data: 0.0122 max
mem: 3407
Epoch: [0] [40/60] eta: 0:00:16 lr: 0.003476 loss: 0.5885 (1.6333)
loss_classifier: 0.0888 (0.2183) loss_box_reg: 0.1846 (0.2099)
loss_mask: 0.2525 (1.1802) loss_objectness: 0.0089 (0.0180)
loss_rpn_box_reg: 0.0080 (0.0069) time: 0.5507 data: 0.0108 max
mem: 3407
Epoch: [0] [50/60] eta: 0:00:07 lr: 0.004323 loss: 0.5317 (1.4095)
loss_classifier: 0.0639 (0.1870) loss_box_reg: 0.2042 (0.2095)
loss_mask: 0.2147 (0.9905) loss_objectness: 0.0048 (0.0155)
loss_rpn_box_reg: 0.0081 (0.0071) time: 0.5623 data: 0.0119 max
mem: 3407
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.4507 (1.2689)
loss_classifier: 0.0513 (0.1673) loss_box_reg: 0.1874 (0.2083)
loss_mask: 0.1876 (0.8724) loss_objectness: 0.0034 (0.0137)
loss_rpn_box_reg: 0.0068 (0.0072) time: 0.5444 data: 0.0113 max
mem: 3407
Epoch: [0] Total time: 0:00:44 (0.7407 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:22 model_time: 0.2382 (0.2382)
evaluator_time: 0.0101 (0.0101) time: 0.4453 data: 0.1947 max mem:
3407
Test: [49/50] eta: 0:00:00 model_time: 0.1134 (0.1350)
evaluator_time: 0.0145 (0.0211) time: 0.1478 data: 0.0039 max mem:
3407
Test: Total time: 0:00:08 (0.1701 s / it)
Averaged stats: model_time: 0.1134 (0.1350) evaluator_time: 0.0145
(0.0211)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.597
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.966
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.661
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.438
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.341
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.618
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.231
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.667
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.672
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.520
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.600
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.683
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.665
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.966
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.841
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.374
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.276
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.689
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.250
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.706
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.710
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.480
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.678
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.722
Epoch: [1] [ 0/60] eta: 0:00:47 lr: 0.005000 loss: 0.2926 (0.2926)
loss_classifier: 0.0359 (0.0359) loss_box_reg: 0.0899 (0.0899)
loss_mask: 0.1515 (0.1515) loss_objectness: 0.0130 (0.0130)
loss_rpn_box_reg: 0.0023 (0.0023) time: 0.7908 data: 0.2946 max
mem: 3407
Epoch: [1] [10/60] eta: 0:00:27 lr: 0.005000 loss: 0.2926 (0.3008)
loss_classifier: 0.0359 (0.0350) loss_box_reg: 0.0916 (0.1098)
loss_mask: 0.1515 (0.1481) loss_objectness: 0.0043 (0.0041)

```



```
loss_rpn_box_reg: 0.0038 (0.0037)  time: 0.5599  data: 0.0327  max
mem: 3407
Epoch: [1]  [20/60]  eta: 0:00:22  lr: 0.005000  loss: 0.2810 (0.2983)
loss_classifier: 0.0300 (0.0350)  loss_box_reg: 0.0915 (0.1056)
loss_mask: 0.1423 (0.1489)  loss_objectness: 0.0016 (0.0036)
loss_rpn_box_reg: 0.0039 (0.0052)  time: 0.5504  data: 0.0087  max
mem: 3407
Epoch: [1]  [30/60]  eta: 0:00:17  lr: 0.005000  loss: 0.2944 (0.3077)
loss_classifier: 0.0388 (0.0381)  loss_box_reg: 0.0845 (0.1094)
loss_mask: 0.1423 (0.1515)  loss_objectness: 0.0009 (0.0029)
loss_rpn_box_reg: 0.0049 (0.0058)  time: 0.5785  data: 0.0103  max
mem: 3407
Epoch: [1]  [40/60]  eta: 0:00:11  lr: 0.005000  loss: 0.2944 (0.3061)
loss_classifier: 0.0428 (0.0379)  loss_box_reg: 0.0987 (0.1085)
loss_mask: 0.1425 (0.1514)  loss_objectness: 0.0008 (0.0028)
loss_rpn_box_reg: 0.0046 (0.0054)  time: 0.5801  data: 0.0086  max
mem: 3407
Epoch: [1]  [50/60]  eta: 0:00:05  lr: 0.005000  loss: 0.2822 (0.3015)
loss_classifier: 0.0318 (0.0366)  loss_box_reg: 0.0833 (0.1049)
loss_mask: 0.1502 (0.1524)  loss_objectness: 0.0008 (0.0025)
loss_rpn_box_reg: 0.0036 (0.0051)  time: 0.5824  data: 0.0094  max
mem: 3407
Epoch: [1]  [59/60]  eta: 0:00:00  lr: 0.005000  loss: 0.3220 (0.3023)
loss_classifier: 0.0328 (0.0373)  loss_box_reg: 0.0833 (0.1041)
loss_mask: 0.1563 (0.1536)  loss_objectness: 0.0007 (0.0024)
loss_rpn_box_reg: 0.0036 (0.0050)  time: 0.5964  data: 0.0091  max
mem: 3407
Epoch: [1] Total time: 0:00:34 (0.5827 s / it)
creating index...
index created!
Test:  [ 0/50]  eta: 0:00:29  model_time: 0.2382 (0.2382)
evaluator_time: 0.0065 (0.0065)  time: 0.5801  data: 0.3342  max mem:
3407
Test:  [49/50]  eta: 0:00:00  model_time: 0.1005 (0.1145)
evaluator_time: 0.0038 (0.0081)  time: 0.1161  data: 0.0036  max mem:
3407
Test: Total time: 0:00:06 (0.1400 s / it)
Averaged stats: model_time: 0.1005 (0.1145)  evaluator_time: 0.0038
(0.0081)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.673
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.970
Average Precision  (AP) @[ IoU=0.75      | area=  all |
```

```

maxDets=100 ] = 0.867
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.444
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.483
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.691
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.265
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.731
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.731
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.460
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.733
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.741
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.695
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.970
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.879
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.404
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.326
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.718
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.272
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.735
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.735
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.722
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.749
Epoch: [2] [ 0/60] eta: 0:00:54 lr: 0.005000 loss: 0.2531 (0.2531)
loss_classifier: 0.0354 (0.0354) loss_box_reg: 0.0526 (0.0526)
loss_mask: 0.1615 (0.1615) loss_objectness: 0.0009 (0.0009)
loss_rpn_box_reg: 0.0027 (0.0027) time: 0.9139 data: 0.3393 max
mem: 3407

```

```
Epoch: [2] [10/60] eta: 0:00:31 lr: 0.005000 loss: 0.2206 (0.2512)
loss_classifier: 0.0274 (0.0311) loss_box_reg: 0.0551 (0.0679)
loss_mask: 0.1279 (0.1456) loss_objectness: 0.0010 (0.0027)
loss_rpn_box_reg: 0.0039 (0.0039) time: 0.6235 data: 0.0387 max
mem: 3407
Epoch: [2] [20/60] eta: 0:00:23 lr: 0.005000 loss: 0.2161 (0.2409)
loss_classifier: 0.0267 (0.0304) loss_box_reg: 0.0551 (0.0619)
loss_mask: 0.1256 (0.1421) loss_objectness: 0.0012 (0.0026)
loss_rpn_box_reg: 0.0039 (0.0038) time: 0.5759 data: 0.0092 max
mem: 3407
Epoch: [2] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2111 (0.2378)
loss_classifier: 0.0245 (0.0302) loss_box_reg: 0.0492 (0.0612)
loss_mask: 0.1246 (0.1403) loss_objectness: 0.0009 (0.0022)
loss_rpn_box_reg: 0.0031 (0.0039) time: 0.5549 data: 0.0084 max
mem: 3407
Epoch: [2] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.2152 (0.2376)
loss_classifier: 0.0258 (0.0303) loss_box_reg: 0.0588 (0.0628)
loss_mask: 0.1315 (0.1385) loss_objectness: 0.0008 (0.0019)
loss_rpn_box_reg: 0.0036 (0.0040) time: 0.5718 data: 0.0087 max
mem: 3407
Epoch: [2] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2153 (0.2311)
loss_classifier: 0.0287 (0.0294) loss_box_reg: 0.0609 (0.0618)
loss_mask: 0.1174 (0.1343) loss_objectness: 0.0007 (0.0018)
loss_rpn_box_reg: 0.0036 (0.0038) time: 0.6102 data: 0.0094 max
mem: 3407
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2048 (0.2284)
loss_classifier: 0.0254 (0.0288) loss_box_reg: 0.0579 (0.0608)
loss_mask: 0.1157 (0.1336) loss_objectness: 0.0005 (0.0016)
loss_rpn_box_reg: 0.0027 (0.0037) time: 0.5979 data: 0.0079 max
mem: 3407
Epoch: [2] Total time: 0:00:35 (0.5909 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:20 model_time: 0.1705 (0.1705)
evaluator_time: 0.0044 (0.0044) time: 0.4113 data: 0.2351 max mem:
3407
Test: [49/50] eta: 0:00:00 model_time: 0.0984 (0.1079)
evaluator_time: 0.0031 (0.0056) time: 0.1135 data: 0.0037 max mem:
3407
Test: Total time: 0:00:06 (0.1261 s / it)
Averaged stats: model_time: 0.0984 (0.1079) evaluator_time: 0.0031
(0.0056)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.741
```

```

Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.913
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.308
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.515
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.765
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.296
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.793
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.793
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.380
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.733
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.813
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.722
Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.967
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.889
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.295
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.342
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.747
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.285
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.772
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.774
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.380
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.744
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.791
Epoch: [3] [ 0/60] eta: 0:00:50  lr: 0.000500  loss: 0.1559 (0.1559)
loss_classifier: 0.0216 (0.0216)  loss_box_reg: 0.0263 (0.0263)

```

```
loss_mask: 0.1068 (0.1068) loss_objectness: 0.0001 (0.0001)
loss_rpn_box_reg: 0.0011 (0.0011) time: 0.8373 data: 0.2845 max
mem: 3407
Epoch: [3] [10/60] eta: 0:00:31 lr: 0.000500 loss: 0.1706 (0.1884)
loss_classifier: 0.0238 (0.0223) loss_box_reg: 0.0357 (0.0448)
loss_mask: 0.1068 (0.1167) loss_objectness: 0.0005 (0.0016)
loss_rpn_box_reg: 0.0021 (0.0030) time: 0.6253 data: 0.0351 max
mem: 3407
Epoch: [3] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1706 (0.1887)
loss_classifier: 0.0238 (0.0246) loss_box_reg: 0.0357 (0.0439)
loss_mask: 0.1093 (0.1158) loss_objectness: 0.0004 (0.0013)
loss_rpn_box_reg: 0.0020 (0.0030) time: 0.5789 data: 0.0088 max
mem: 3407
Epoch: [3] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1689 (0.1857)
loss_classifier: 0.0207 (0.0242) loss_box_reg: 0.0273 (0.0413)
loss_mask: 0.1099 (0.1165) loss_objectness: 0.0004 (0.0011)
loss_rpn_box_reg: 0.0020 (0.0027) time: 0.5743 data: 0.0090 max
mem: 3407
Epoch: [3] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1799 (0.1888)
loss_classifier: 0.0198 (0.0242) loss_box_reg: 0.0323 (0.0408)
loss_mask: 0.1154 (0.1199) loss_objectness: 0.0004 (0.0012)
loss_rpn_box_reg: 0.0025 (0.0027) time: 0.5977 data: 0.0102 max
mem: 3407
Epoch: [3] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1842 (0.1881)
loss_classifier: 0.0198 (0.0232) loss_box_reg: 0.0326 (0.0406)
loss_mask: 0.1230 (0.1204) loss_objectness: 0.0005 (0.0011)
loss_rpn_box_reg: 0.0026 (0.0028) time: 0.5926 data: 0.0091 max
mem: 3407
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1908 (0.1891)
loss_classifier: 0.0205 (0.0234) loss_box_reg: 0.0351 (0.0413)
loss_mask: 0.1183 (0.1204) loss_objectness: 0.0005 (0.0011)
loss_rpn_box_reg: 0.0024 (0.0028) time: 0.5856 data: 0.0084 max
mem: 3407
Epoch: [3] Total time: 0:00:35 (0.5951 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:20 model_time: 0.1936 (0.1936)
evaluator_time: 0.0046 (0.0046) time: 0.4008 data: 0.2013 max mem:
3407
Test: [49/50] eta: 0:00:00 model_time: 0.1010 (0.1097)
evaluator_time: 0.0041 (0.0061) time: 0.1166 data: 0.0038 max mem:
3407
Test: Total time: 0:00:06 (0.1297 s / it)
Averaged stats: model_time: 0.1010 (0.1097) evaluator_time: 0.0041
(0.0061)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
```

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.791

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.973

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.943

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.334

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.569

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.817

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=1 ] = 0.313

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=10 ] = 0.837

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.837

Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.400

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.767

Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.858

IoU metric: segm

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.729

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.969

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.895

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.285

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.372

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.755

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=1 ] = 0.287

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=10 ] = 0.780

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.780

Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.380

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.733

Average Recall (AR) @[ IoU=0.50:0.95 | area= large |

```
maxDets=100 ] = 0.798
Epoch: [4] [ 0/60] eta: 0:01:14 lr: 0.000500 loss: 0.2000 (0.2000)
loss_classifier: 0.0283 (0.0283) loss_box_reg: 0.0574 (0.0574)
loss_mask: 0.1124 (0.1124) loss_objectness: 0.0004 (0.0004)
loss_rpn_box_reg: 0.0015 (0.0015) time: 1.2390 data: 0.4964 max
mem: 3407
Epoch: [4] [10/60] eta: 0:00:32 lr: 0.000500 loss: 0.1837 (0.1751)
loss_classifier: 0.0210 (0.0233) loss_box_reg: 0.0352 (0.0367)
loss_mask: 0.1124 (0.1119) loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0022 (0.0023) time: 0.6400 data: 0.0528 max
mem: 3407
Epoch: [4] [20/60] eta: 0:00:24 lr: 0.000500 loss: 0.1549 (0.1658)
loss_classifier: 0.0165 (0.0201) loss_box_reg: 0.0288 (0.0326)
loss_mask: 0.1073 (0.1103) loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0022 (0.0022) time: 0.5862 data: 0.0086 max
mem: 3410
Epoch: [4] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.1382 (0.1709)
loss_classifier: 0.0147 (0.0215) loss_box_reg: 0.0230 (0.0342)
loss_mask: 0.1014 (0.1120) loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0019 (0.0025) time: 0.5929 data: 0.0090 max
mem: 3410
Epoch: [4] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1879 (0.1780)
loss_classifier: 0.0206 (0.0223) loss_box_reg: 0.0416 (0.0374)
loss_mask: 0.1151 (0.1150) loss_objectness: 0.0006 (0.0007)
loss_rpn_box_reg: 0.0025 (0.0026) time: 0.5954 data: 0.0085 max
mem: 3410
Epoch: [4] [50/60] eta: 0:00:06 lr: 0.000500 loss: 0.1976 (0.1836)
loss_classifier: 0.0208 (0.0232) loss_box_reg: 0.0457 (0.0395)
loss_mask: 0.1255 (0.1173) loss_objectness: 0.0006 (0.0008)
loss_rpn_box_reg: 0.0025 (0.0027) time: 0.6122 data: 0.0094 max
mem: 3410
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1911 (0.1855)
loss_classifier: 0.0293 (0.0234) loss_box_reg: 0.0457 (0.0398)
loss_mask: 0.1134 (0.1188) loss_objectness: 0.0005 (0.0008)
loss_rpn_box_reg: 0.0022 (0.0027) time: 0.5879 data: 0.0092 max
mem: 3410
Epoch: [4] Total time: 0:00:36 (0.6034 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:21 model_time: 0.1406 (0.1406)
evaluator_time: 0.0039 (0.0039) time: 0.4317 data: 0.2859 max mem:
3410
Test: [49/50] eta: 0:00:00 model_time: 0.1012 (0.1153)
evaluator_time: 0.0039 (0.0086) time: 0.1174 data: 0.0039 max mem:
3410
Test: Total time: 0:00:07 (0.1409 s / it)
Averaged stats: model_time: 0.1012 (0.1153) evaluator_time: 0.0039
(0.0086)
Accumulating evaluation results...
```

```

DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.790
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.921
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.369
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.575
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.816
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.311
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.842
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.842
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.420
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.778
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.862
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.737
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.968
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.902
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.306
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.365
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.763
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.287
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.784
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.786
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |

```



```
maxDets=100 ] = 0.733
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.805
Epoch: [5] [ 0/60] eta: 0:00:54 lr: 0.000500 loss: 0.1570 (0.1570)
loss_classifier: 0.0241 (0.0241) loss_box_reg: 0.0265 (0.0265)
loss_mask: 0.1041 (0.1041) loss_objectness: 0.0005 (0.0005)
loss_rpn_box_reg: 0.0019 (0.0019) time: 0.9105 data: 0.2797 max
mem: 3410
Epoch: [5] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.1600 (0.1809)
loss_classifier: 0.0242 (0.0243) loss_box_reg: 0.0299 (0.0363)
loss_mask: 0.1113 (0.1160) loss_objectness: 0.0010 (0.0020)
loss_rpn_box_reg: 0.0019 (0.0023) time: 0.6086 data: 0.0319 max
mem: 3410
Epoch: [5] [20/60] eta: 0:00:24 lr: 0.000500 loss: 0.1600 (0.1764)
loss_classifier: 0.0226 (0.0246) loss_box_reg: 0.0269 (0.0348)
loss_mask: 0.1113 (0.1127) loss_objectness: 0.0006 (0.0015)
loss_rpn_box_reg: 0.0019 (0.0028) time: 0.6059 data: 0.0084 max
mem: 3410
Epoch: [5] [30/60] eta: 0:00:18 lr: 0.000500 loss: 0.1522 (0.1732)
loss_classifier: 0.0186 (0.0229) loss_box_reg: 0.0269 (0.0349)
loss_mask: 0.1044 (0.1116) loss_objectness: 0.0004 (0.0012)
loss_rpn_box_reg: 0.0016 (0.0025) time: 0.6002 data: 0.0092 max
mem: 3410
Epoch: [5] [40/60] eta: 0:00:12 lr: 0.000500 loss: 0.1547 (0.1733)
loss_classifier: 0.0197 (0.0228) loss_box_reg: 0.0320 (0.0343)
loss_mask: 0.1059 (0.1126) loss_objectness: 0.0003 (0.0012)
loss_rpn_box_reg: 0.0015 (0.0023) time: 0.5807 data: 0.0090 max
mem: 3779
Epoch: [5] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1674 (0.1724)
loss_classifier: 0.0222 (0.0229) loss_box_reg: 0.0320 (0.0342)
loss_mask: 0.1059 (0.1118) loss_objectness: 0.0003 (0.0012)
loss_rpn_box_reg: 0.0015 (0.0023) time: 0.5845 data: 0.0097 max
mem: 3779
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1864 (0.1783)
loss_classifier: 0.0237 (0.0231) loss_box_reg: 0.0342 (0.0361)
loss_mask: 0.1142 (0.1155) loss_objectness: 0.0004 (0.0011)
loss_rpn_box_reg: 0.0021 (0.0025) time: 0.5814 data: 0.0085 max
mem: 3779
Epoch: [5] Total time: 0:00:35 (0.5992 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:19 model_time: 0.1444 (0.1444)
evaluator_time: 0.0061 (0.0061) time: 0.3975 data: 0.2457 max mem:
3779
Test: [49/50] eta: 0:00:00 model_time: 0.1005 (0.1083)
evaluator_time: 0.0030 (0.0054) time: 0.1134 data: 0.0036 max mem:
3779
Test: Total time: 0:00:06 (0.1264 s / it)
Averaged stats: model_time: 0.1005 (0.1083) evaluator_time: 0.0030
```

```

(0.0054)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.793
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.930
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.339
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.589
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.819
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.315
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.844
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.844
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.420
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.864
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.737
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.916
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.308
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.347
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.760
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.288
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.786
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.787

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.420
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.767
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.802
Epoch: [6] [ 0/60] eta: 0:01:00 lr: 0.000050 loss: 0.2225 (0.2225)
loss_classifier: 0.0314 (0.0314) loss_box_reg: 0.0443 (0.0443)
loss_mask: 0.1423 (0.1423) loss_objectness: 0.0002 (0.0002)
loss_rpn_box_reg: 0.0043 (0.0043) time: 1.0073 data: 0.3170 max
mem: 3779
Epoch: [6] [10/60] eta: 0:00:33 lr: 0.000050 loss: 0.1692 (0.1704)
loss_classifier: 0.0220 (0.0220) loss_box_reg: 0.0307 (0.0314)
loss_mask: 0.1047 (0.1130) loss_objectness: 0.0005 (0.0014)
loss_rpn_box_reg: 0.0022 (0.0026) time: 0.6642 data: 0.0387 max
mem: 3779
Epoch: [6] [20/60] eta: 0:00:24 lr: 0.000050 loss: 0.1605 (0.1715)
loss_classifier: 0.0202 (0.0224) loss_box_reg: 0.0281 (0.0313)
loss_mask: 0.1100 (0.1145) loss_objectness: 0.0005 (0.0012)
loss_rpn_box_reg: 0.0017 (0.0021) time: 0.5981 data: 0.0097 max
mem: 3779
Epoch: [6] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1690 (0.1744)
loss_classifier: 0.0202 (0.0232) loss_box_reg: 0.0281 (0.0323)
loss_mask: 0.1133 (0.1156) loss_objectness: 0.0003 (0.0012)
loss_rpn_box_reg: 0.0015 (0.0020) time: 0.5603 data: 0.0089 max
mem: 3779
Epoch: [6] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1902 (0.1749)
loss_classifier: 0.0248 (0.0236) loss_box_reg: 0.0345 (0.0336)
loss_mask: 0.1122 (0.1144) loss_objectness: 0.0005 (0.0011)
loss_rpn_box_reg: 0.0019 (0.0022) time: 0.5786 data: 0.0098 max
mem: 3779
Epoch: [6] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1706 (0.1746)
loss_classifier: 0.0248 (0.0231) loss_box_reg: 0.0345 (0.0336)
loss_mask: 0.1102 (0.1146) loss_objectness: 0.0005 (0.0010)
loss_rpn_box_reg: 0.0027 (0.0023) time: 0.5921 data: 0.0095 max
mem: 3779
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1706 (0.1763)
loss_classifier: 0.0233 (0.0228) loss_box_reg: 0.0322 (0.0339)
loss_mask: 0.1174 (0.1163) loss_objectness: 0.0004 (0.0010)
loss_rpn_box_reg: 0.0029 (0.0024) time: 0.6022 data: 0.0081 max
mem: 3779
Epoch: [6] Total time: 0:00:36 (0.6023 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:20 model_time: 0.1700 (0.1700)
evaluator_time: 0.0040 (0.0040) time: 0.4024 data: 0.2271 max mem:
3779
Test: [49/50] eta: 0:00:00 model_time: 0.1016 (0.1096)
evaluator_time: 0.0041 (0.0057) time: 0.1160 data: 0.0039 max mem:

```

3779

Test: Total time: 0:00:06 (0.1293 s / it)

Averaged stats: model\_time: 0.1016 (0.1096) evaluator\_time: 0.0041 (0.0057)

Accumulating evaluation results...

DONE (t=0.03s).

Accumulating evaluation results...

DONE (t=0.02s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.792

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.972

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.930

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.334

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.589

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.818

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=1 ] = 0.315

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=10 ] = 0.842

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.842

Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.400

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.789

Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.863

IoU metric: segm

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.734

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.968

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.908

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.299

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.339

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.759

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=1 ] = 0.289

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=

```
10 ] = 0.781
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.784
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.744
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.802
Epoch: [7] [ 0/60] eta: 0:01:10  lr: 0.000050  loss: 0.1581 (0.1581)
loss_classifier: 0.0195 (0.0195)  loss_box_reg: 0.0305 (0.0305)
loss_mask: 0.1040 (0.1040)  loss_objectness: 0.0011 (0.0011)
loss_rpn_box_reg: 0.0030 (0.0030)  time: 1.1739  data: 0.5194  max
mem: 3779
Epoch: [7] [10/60] eta: 0:00:32  lr: 0.000050  loss: 0.1719 (0.1826)
loss_classifier: 0.0195 (0.0235)  loss_box_reg: 0.0384 (0.0371)
loss_mask: 0.1124 (0.1185)  loss_objectness: 0.0004 (0.0006)
loss_rpn_box_reg: 0.0028 (0.0029)  time: 0.6466  data: 0.0533  max
mem: 3779
Epoch: [7] [20/60] eta: 0:00:24  lr: 0.000050  loss: 0.1719 (0.1936)
loss_classifier: 0.0225 (0.0260)  loss_box_reg: 0.0384 (0.0400)
loss_mask: 0.1154 (0.1236)  loss_objectness: 0.0004 (0.0011)
loss_rpn_box_reg: 0.0025 (0.0029)  time: 0.5911  data: 0.0080  max
mem: 3779
Epoch: [7] [30/60] eta: 0:00:18  lr: 0.000050  loss: 0.1718 (0.1838)
loss_classifier: 0.0200 (0.0241)  loss_box_reg: 0.0291 (0.0366)
loss_mask: 0.1131 (0.1198)  loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0022 (0.0026)  time: 0.5773  data: 0.0096  max
mem: 3779
Epoch: [7] [40/60] eta: 0:00:11  lr: 0.000050  loss: 0.1669 (0.1785)
loss_classifier: 0.0179 (0.0226)  loss_box_reg: 0.0289 (0.0351)
loss_mask: 0.1117 (0.1175)  loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0017 (0.0025)  time: 0.5661  data: 0.0088  max
mem: 3779
Epoch: [7] [50/60] eta: 0:00:05  lr: 0.000050  loss: 0.1603 (0.1758)
loss_classifier: 0.0179 (0.0221)  loss_box_reg: 0.0281 (0.0341)
loss_mask: 0.1034 (0.1162)  loss_objectness: 0.0004 (0.0009)
loss_rpn_box_reg: 0.0021 (0.0024)  time: 0.5967  data: 0.0105  max
mem: 3779
Epoch: [7] [59/60] eta: 0:00:00  lr: 0.000050  loss: 0.1628 (0.1761)
loss_classifier: 0.0193 (0.0229)  loss_box_reg: 0.0281 (0.0345)
loss_mask: 0.1034 (0.1154)  loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0018 (0.0024)  time: 0.5929  data: 0.0101  max
mem: 3779
Epoch: [7] Total time: 0:00:35 (0.5961 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:20  model_time: 0.1477 (0.1477)
evaluator_time: 0.0046 (0.0046)  time: 0.4061  data: 0.2524  max mem:
```

```

3779
Test: [49/50] eta: 0:00:00 model_time: 0.1069 (0.1163)
evaluator_time: 0.0037 (0.0077) time: 0.1259 data: 0.0062 max mem:
3779
Test: Total time: 0:00:07 (0.1406 s / it)
Averaged stats: model_time: 0.1069 (0.1163) evaluator_time: 0.0037
(0.0077)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.792
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.922
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.334
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.589
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.817
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.315
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.841
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.841
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.862
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.735
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.908
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.308
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.333
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.761

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.289
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.783
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.785
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.420
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.733
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.803
Epoch: [8] [ 0/60] eta: 0:01:01  lr: 0.000050  loss: 0.1972 (0.1972)
loss_classifier: 0.0473 (0.0473)  loss_box_reg: 0.0457 (0.0457)
loss_mask: 0.1003 (0.1003)  loss_objectness: 0.0012 (0.0012)
loss_rpn_box_reg: 0.0027 (0.0027)  time: 1.0282  data: 0.3046  max
mem: 3779
Epoch: [8] [10/60] eta: 0:00:31  lr: 0.000050  loss: 0.1721 (0.1809)
loss_classifier: 0.0237 (0.0249)  loss_box_reg: 0.0286 (0.0378)
loss_mask: 0.1126 (0.1152)  loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0020 (0.0022)  time: 0.6214  data: 0.0346  max
mem: 3779
Epoch: [8] [20/60] eta: 0:00:24  lr: 0.000050  loss: 0.1669 (0.1819)
loss_classifier: 0.0207 (0.0244)  loss_box_reg: 0.0286 (0.0392)
loss_mask: 0.1121 (0.1152)  loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0019 (0.0025)  time: 0.5878  data: 0.0098  max
mem: 3779
Epoch: [8] [30/60] eta: 0:00:17  lr: 0.000050  loss: 0.1732 (0.1791)
loss_classifier: 0.0221 (0.0240)  loss_box_reg: 0.0316 (0.0371)
loss_mask: 0.1121 (0.1147)  loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0019 (0.0024)  time: 0.5788  data: 0.0101  max
mem: 3779
Epoch: [8] [40/60] eta: 0:00:11  lr: 0.000050  loss: 0.1644 (0.1758)
loss_classifier: 0.0187 (0.0230)  loss_box_reg: 0.0300 (0.0364)
loss_mask: 0.1056 (0.1133)  loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0013 (0.0023)  time: 0.5661  data: 0.0101  max
mem: 3779
Epoch: [8] [50/60] eta: 0:00:05  lr: 0.000050  loss: 0.1460 (0.1709)
loss_classifier: 0.0148 (0.0217)  loss_box_reg: 0.0224 (0.0347)
loss_mask: 0.1005 (0.1115)  loss_objectness: 0.0003 (0.0007)
loss_rpn_box_reg: 0.0020 (0.0023)  time: 0.5697  data: 0.0099  max
mem: 3779
Epoch: [8] [59/60] eta: 0:00:00  lr: 0.000050  loss: 0.1596 (0.1747)
loss_classifier: 0.0173 (0.0223)  loss_box_reg: 0.0295 (0.0356)
loss_mask: 0.1069 (0.1137)  loss_objectness: 0.0003 (0.0006)
loss_rpn_box_reg: 0.0024 (0.0024)  time: 0.5978  data: 0.0081  max
mem: 3779
Epoch: [8] Total time: 0:00:35 (0.5921 s / it)
creating index...

```

```

index created!
Test: [ 0/50] eta: 0:00:29 model_time: 0.2405 (0.2405)
evaluator_time: 0.0040 (0.0040) time: 0.5880 data: 0.3419 max mem:
3779
Test: [49/50] eta: 0:00:00 model_time: 0.1009 (0.1111)
evaluator_time: 0.0030 (0.0055) time: 0.1148 data: 0.0037 max mem:
3779
Test: Total time: 0:00:06 (0.1317 s / it)
Averaged stats: model_time: 0.1009 (0.1111) evaluator_time: 0.0030
(0.0055)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.792
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.922
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.334
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.589
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.817
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.316
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.841
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.841
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.862
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.735
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.908
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.308
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |

```



```
maxDets=100 ] = 0.336
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.760
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.288
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.785
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.420
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.744
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.802
Epoch: [9] [ 0/60] eta: 0:00:56 lr: 0.000005 loss: 0.3186 (0.3186)
loss_classifier: 0.0453 (0.0453) loss_box_reg: 0.1030 (0.1030)
loss_mask: 0.1553 (0.1553) loss_objectness: 0.0088 (0.0088)
loss_rpn_box_reg: 0.0061 (0.0061) time: 0.9412 data: 0.3493 max
mem: 3779
Epoch: [9] [10/60] eta: 0:00:29 lr: 0.000005 loss: 0.1587 (0.1720)
loss_classifier: 0.0184 (0.0197) loss_box_reg: 0.0285 (0.0338)
loss_mask: 0.1106 (0.1150) loss_objectness: 0.0003 (0.0013)
loss_rpn_box_reg: 0.0019 (0.0023) time: 0.5939 data: 0.0393 max
mem: 3779
Epoch: [9] [20/60] eta: 0:00:23 lr: 0.000005 loss: 0.1529 (0.1661)
loss_classifier: 0.0172 (0.0202) loss_box_reg: 0.0282 (0.0323)
loss_mask: 0.1032 (0.1105) loss_objectness: 0.0003 (0.0009)
loss_rpn_box_reg: 0.0019 (0.0021) time: 0.5664 data: 0.0081 max
mem: 3779
Epoch: [9] [30/60] eta: 0:00:17 lr: 0.000005 loss: 0.1641 (0.1727)
loss_classifier: 0.0190 (0.0211) loss_box_reg: 0.0297 (0.0341)
loss_mask: 0.1054 (0.1145) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0019 (0.0021) time: 0.5646 data: 0.0082 max
mem: 3779
Epoch: [9] [40/60] eta: 0:00:11 lr: 0.000005 loss: 0.1668 (0.1693)
loss_classifier: 0.0190 (0.0207) loss_box_reg: 0.0282 (0.0328)
loss_mask: 0.1064 (0.1127) loss_objectness: 0.0004 (0.0008)
loss_rpn_box_reg: 0.0019 (0.0022) time: 0.5789 data: 0.0086 max
mem: 3779
Epoch: [9] [50/60] eta: 0:00:05 lr: 0.000005 loss: 0.1668 (0.1765)
loss_classifier: 0.0227 (0.0227) loss_box_reg: 0.0283 (0.0349)
loss_mask: 0.1100 (0.1158) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0019 (0.0025) time: 0.6191 data: 0.0088 max
mem: 3779
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1646 (0.1758)
loss_classifier: 0.0197 (0.0221) loss_box_reg: 0.0283 (0.0343)
loss_mask: 0.1131 (0.1162) loss_objectness: 0.0003 (0.0008)
loss_rpn_box_reg: 0.0018 (0.0024) time: 0.6305 data: 0.0085 max
```

```
mem: 3779
Epoch: [9] Total time: 0:00:35 (0.5994 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:21 model_time: 0.1965 (0.1965)
evaluator_time: 0.0040 (0.0040) time: 0.4318 data: 0.2301 max mem:
3779
Test: [49/50] eta: 0:00:00 model_time: 0.1138 (0.1178)
evaluator_time: 0.0051 (0.0068) time: 0.1340 data: 0.0062 max mem:
3779
Test: Total time: 0:00:06 (0.1386 s / it)
Averaged stats: model_time: 0.1138 (0.1178) evaluator_time: 0.0051
(0.0068)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.791
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.922
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.334
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.589
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.817
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.316
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.840
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.840
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.860
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.734
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.972
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.908
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.299
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.336
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.760
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.288
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.784
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.744
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.802
That's it!

```

### Comments on training log:

In the above training log, note the last batch of the 10th epoch result.

```

Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1646 (0.1758) loss_classifier: 0.0197
(0.0221) loss_box_reg: 0.0283 (0.0343) loss_mask: 0.1131 (0.1162) loss_objectness: 0.0003
(0.0008) loss_rpn_box_reg: 0.0018 (0.0024) time: 0.6305 data: 0.0085 max mem: 3779

```

Here, the loss value in paranthesis represents the cumulative loss over the entire epoch upto that point( here its the last batch so its for the entire epoch ) using the weights after the completion of the 10th epoch.

Similarly for loss\_classifier, loss\_mask, loss\_objectness, loss\_rpn\_box\_reg

These results can be used for comparing the model performance asked in Q5B)

### Testing of finetuned model on Beatles\_Abbey\_Road Test image: (Method 1)

```

import matplotlib.pyplot as plt
import cv2
from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

# Read an image from a specified path
image = read_image("/content/sample_data/Beatles_-_Abbey_Road.jpeg")

# Create an output image to visualize the results
output_image = image

# Obtain an evaluation transformation with 'train=False'
eval_transform = get_transform(train=False)

```

```

# Set the model in evaluation mode
model.eval()

with torch.no_grad():
    x = eval_transform(image)

    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)

    # Make predictions using the model
    predictions = model([x, ])
    pred = predictions[0]

# Normalize and convert the image to 8-bit integers (uint8)
image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]

# Filter predictions based on confidence scores (only keep scores >
0.65)
# mask refers to binary-mask ie true, false of predictions above
confidence( not meaning the mask displayed in the image)
mask = pred["scores"] > 0.65
filtered_pred = {key: value[mask] for key, value in pred.items()}

#Obtaining labels, boxes and masks for filtered predictions
filtered_labels = [f"ped: {score:.3f}" for score in
filtered_pred["scores"]]
filtered_boxes = filtered_pred["boxes"].long()
masks = (filtered_pred["masks"] > 0.7).squeeze(1)

#output image having the filtered prediction masks now
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

# Convert to NumPy array
output_image = output_image.permute(1, 2, 0).cpu().numpy().copy()


#Drawing the boxes, labels using cv2
for label, box in zip(filtered_labels, filtered_boxes):
    x_1, y_1, x_2, y_2 = [coord.item() for coord in box]
    output_image = cv2.rectangle(output_image, (x_1, y_1), (x_2, y_2),
(255, 0, 0), 2)
    output_image = cv2.putText(output_image, label, (x_1, y_1 - 10),
cv2.FONT_HERSHEY_SIMPLEX, 0.3, (255, 0, 0), 1)

#Plotting the final image
plt.figure(figsize=(10, 10))

```

```
plt.imshow(output_image)
```

```
<matplotlib.image.AxesImage at 0x7a4c54537cd0>
```



```
print(pred["scores"])
```

```
tensor([0.9806, 0.9776, 0.9639, 0.9368, 0.1338, 0.1002, 0.0907],  
        device='cuda:0')
```

**Comments on testing (method 1) result:**

In the above method for testing, predictions were inferred from the model. Each prediction dictionary consisted of labels, boxes, masks, confidence scores as keys.

Using confidence score threshold of 0.7, predictions in the list were filtered and the corresponding boxes, labels, masks for the the filtered predictions were outputed. In this case, out of all prediction scores (as printed above) four of them crossed the threshold.

### Testing of finetuned model on Beatles\_Abbey\_Road Test image: (Method 2)

```
import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

# Read an image from a specified path
image = read_image("/content/sample_data/Beatles_-_Abbey_Road.jpeg")

# Create an output image to visualize the results
output_image = image

# Obtain an evaluation transformation with 'train=False'
eval_transform = get_transform(train=False)

# Set the model in evaluation mode
model.eval()

with torch.no_grad():
    x = eval_transform(image)

    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)

    # Make predictions using the model
    predictions = model([x, ])
    pred = predictions[0]

# Normalize and convert the image to 8-bit integers (uint8)
image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]

#Filtering masks based on confidence
masks = (pred["masks"] > 0.7).squeeze(1)


#output image having the filtered prediction masks now
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

#Plotting the final output image
plt.figure(figsize=(12, 12))
```



```
plt.imshow(output_image.permute(1, 2, 0))
```

```
<matplotlib.image.AxesImage at 0x7a4c54332e00>
```



### Comments on testing (method 2) result:

In the above method for testing, predictions were inferred from the model.

From each prediction, the corresponding masks whose values exceeded the set threshold were used for the output image.

Prediction dictionaries with low scores like 0.1 etc might have some values in their masks as 0.8 etc (above the threshold). These masks are included in this method but discarded in the previous method.

## 2 - Modifying the model to add a different backbone - Mobilenet ( Option 2)

```
import torchvision
from torchvision.models.detection import FasterRCNN
from torchvision.models.detection.rpn import AnchorGenerator
from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor
from torchvision.models.detection import MaskRCNN

def get_backbone_model_instance_segmentation(num_classes):
    # Load a pre-trained model for classification and return only the features
    backbone =
torchvision.models.mobilenet_v2(weights="DEFAULT").features
    # Set the number of output channels in the backbone to 1280
    backbone.out_channels = 1280

    # Define the anchor generator with desired anchor sizes and aspect ratios
    anchor_generator = AnchorGenerator(
        sizes=((32, 64, 128, 256, 512)),
        aspect_ratios=((0.5, 1.0, 2.0)),
    )

    # Define the feature maps to use for region of interest cropping and resizing
    roi_pooler = torchvision.ops.MultiScaleRoIAlign(
        featmap_names=['0'],
        output_size=7,
        sampling_ratio=2,
    )

    mask_roi_pooler =
torchvision.ops.MultiScaleRoIAlign(featmap_names=['0'],
output_size=14,
sampling_ratio=2)

    # Create a Mask R-CNN model with the custom backbone
    model = MaskRCNN(
        backbone,
        num_classes=num_classes,
        rpn_anchor_generator=anchor_generator,
        box_roi_pool=roi_pooler,
```



```

        mask_roi_pool=mask_roi_pooler
    )

    return model

```

Q5A) The above is the code for the modified backbone model.

### ***Training and validation for our Modified backbone Model***

```

from engine import train_one_epoch, evaluate

# train on the GPU or on the CPU, if a GPU is not available
device = torch.device('cuda') if torch.cuda.is_available() else
torch.device('cpu')

# our dataset has two classes only - background and person
num_classes = 2
# use our dataset and defined transformations
dataset = PennFudanDataset('/content/drive/MyDrive/PennFudanPed',
    get_transform(train=True))
dataset_test = PennFudanDataset('/content/drive/MyDrive/PennFudanPed',
    get_transform(train=False))

# split the dataset in train and test set
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset,
    batch_size=2,
    shuffle=True,
    num_workers=4,
    collate_fn=utils.collate_fn
)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test,
    batch_size=1,
    shuffle=False,
    num_workers=4,
    collate_fn=utils.collate_fn
)

# get the model using our helper function
model = get_backbone_model_instance_segmentation(num_classes)

```

```

# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(
    params,
    lr=0.005,
    momentum=0.9,
    weight_decay=0.0005
)

# and a learning rate scheduler
lr_scheduler = torch.optim.lr_scheduler.StepLR(
    optimizer,
    step_size=3,
    gamma=0.1
)

# let's train it for 5 epochs
num_epochs = 10

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch,
    print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)

print("That's it!")

/usr/local/lib/python3.10/dist-packages/torch/utils/data/
dataloader.py:557: UserWarning: This DataLoader will create 4 worker
processes in total. Our suggested max number of worker in current
system is 2, which is smaller than what this DataLoader is going to
create. Please be aware that excessive worker creation might get
DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(
Downloading: "https://download.pytorch.org/models/mobilenet_v2-
7ebf99e0.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-
7ebf99e0.pth
100%|██████████| 13.6M/13.6M [00:00<00:00, 27.0MB/s]

Epoch: [0] [ 0/60] eta: 0:01:18 lr: 0.000090 loss: 3.8727 (3.8727)
loss_classifier: 0.6975 (0.6975) loss_box_reg: 0.0829 (0.0829)
loss_mask: 2.2960 (2.2960) loss_objectness: 0.6972 (0.6972)
loss_rpn_box_reg: 0.0992 (0.0992) time: 1.3155 data: 0.4969 max

```

```
mem: 4376
Epoch: [0] [10/60] eta: 0:00:23 lr: 0.000936 loss: 3.5443 (3.4568)
loss_classifier: 0.6548 (0.6223) loss_box_reg: 0.0848 (0.0949)
loss_mask: 2.0635 (2.0079) loss_objectness: 0.6926 (0.6870)
loss_rpn_box_reg: 0.0426 (0.0448) time: 0.4704 data: 0.0509 max
mem: 5168
Epoch: [0] [20/60] eta: 0:00:17 lr: 0.001783 loss: 2.6578 (2.8665)
loss_classifier: 0.4284 (0.4715) loss_box_reg: 0.1107 (0.1281)
loss_mask: 1.3966 (1.5790) loss_objectness: 0.6510 (0.6464)
loss_rpn_box_reg: 0.0311 (0.0416) time: 0.4059 data: 0.0087 max
mem: 5187
Epoch: [0] [30/60] eta: 0:00:13 lr: 0.002629 loss: 1.7916 (2.4441)
loss_classifier: 0.2523 (0.3883) loss_box_reg: 0.1195 (0.1288)
loss_mask: 0.7981 (1.3042) loss_objectness: 0.5453 (0.5848)
loss_rpn_box_reg: 0.0255 (0.0381) time: 0.4202 data: 0.0102 max
mem: 5238
Epoch: [0] [40/60] eta: 0:00:08 lr: 0.003476 loss: 1.4503 (2.2111)
loss_classifier: 0.2523 (0.3698) loss_box_reg: 0.1484 (0.1424)
loss_mask: 0.6806 (1.1420) loss_objectness: 0.3846 (0.5214)
loss_rpn_box_reg: 0.0230 (0.0354) time: 0.4082 data: 0.0088 max
mem: 5238
Epoch: [0] [50/60] eta: 0:00:04 lr: 0.004323 loss: 1.4167 (2.0313)
loss_classifier: 0.2569 (0.3498) loss_box_reg: 0.1516 (0.1513)
loss_mask: 0.5925 (1.0291) loss_objectness: 0.2670 (0.4668)
loss_rpn_box_reg: 0.0260 (0.0344) time: 0.4050 data: 0.0094 max
mem: 5238
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 1.1461 (1.8956)
loss_classifier: 0.2199 (0.3301) loss_box_reg: 0.1576 (0.1543)
loss_mask: 0.5550 (0.9523) loss_objectness: 0.2200 (0.4261)
loss_rpn_box_reg: 0.0255 (0.0329) time: 0.4047 data: 0.0092 max
mem: 5238
Epoch: [0] Total time: 0:00:25 (0.4244 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:29 model_time: 0.2214 (0.2214)
evaluator_time: 0.0480 (0.0480) time: 0.5894 data: 0.3163 max mem:
5238
Test: [49/50] eta: 0:00:00 model_time: 0.1692 (0.1585)
evaluator_time: 0.0416 (0.0388) time: 0.2242 data: 0.0056 max mem:
5238
Test: Total time: 0:00:10 (0.2145 s / it)
Averaged stats: model_time: 0.1692 (0.1585) evaluator_time: 0.0416
(0.0388)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
```

```

maxDets=100 ] = 0.010
Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.035
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.002
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.093
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.018
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.082
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.310
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.329
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.006
Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.031
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.063
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.009
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.057
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.178
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.189
Epoch: [1] [ 0/60] eta: 0:00:50 lr: 0.005000 loss: 1.1749 (1.1749)

```

```
loss_classifier: 0.2219 (0.2219) loss_box_reg: 0.2162 (0.2162)
loss_mask: 0.5386 (0.5386) loss_objectness: 0.1663 (0.1663)
loss_rpn_box_reg: 0.0319 (0.0319) time: 0.8437 data: 0.3308 max
mem: 5238
Epoch: [1] [10/60] eta: 0:00:22 lr: 0.005000 loss: 1.1403 (1.0908)
loss_classifier: 0.2185 (0.2094) loss_box_reg: 0.2050 (0.1862)
loss_mask: 0.5008 (0.5024) loss_objectness: 0.1663 (0.1667)
loss_rpn_box_reg: 0.0230 (0.0261) time: 0.4479 data: 0.0367 max
mem: 5238
Epoch: [1] [20/60] eta: 0:00:17 lr: 0.005000 loss: 1.0351 (1.0603)
loss_classifier: 0.2005 (0.2010) loss_box_reg: 0.1857 (0.1847)
loss_mask: 0.4656 (0.4918) loss_objectness: 0.1493 (0.1560)
loss_rpn_box_reg: 0.0255 (0.0269) time: 0.4136 data: 0.0088 max
mem: 5238
Epoch: [1] [30/60] eta: 0:00:12 lr: 0.005000 loss: 0.9634 (1.0375)
loss_classifier: 0.1728 (0.1897) loss_box_reg: 0.1574 (0.1869)
loss_mask: 0.4656 (0.4854) loss_objectness: 0.1364 (0.1483)
loss_rpn_box_reg: 0.0255 (0.0272) time: 0.4240 data: 0.0105 max
mem: 5238
Epoch: [1] [40/60] eta: 0:00:08 lr: 0.005000 loss: 0.8829 (1.0028)
loss_classifier: 0.1422 (0.1786) loss_box_reg: 0.1574 (0.1845)
loss_mask: 0.4474 (0.4745) loss_objectness: 0.1141 (0.1388)
loss_rpn_box_reg: 0.0228 (0.0263) time: 0.4159 data: 0.0096 max
mem: 5238
Epoch: [1] [50/60] eta: 0:00:04 lr: 0.005000 loss: 0.8409 (0.9674)
loss_classifier: 0.1318 (0.1692) loss_box_reg: 0.1496 (0.1809)
loss_mask: 0.4180 (0.4609) loss_objectness: 0.0969 (0.1306)
loss_rpn_box_reg: 0.0226 (0.0258) time: 0.4066 data: 0.0087 max
mem: 6183
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.7235 (0.9345)
loss_classifier: 0.1133 (0.1602) loss_box_reg: 0.1133 (0.1743)
loss_mask: 0.3708 (0.4501) loss_objectness: 0.0819 (0.1242)
loss_rpn_box_reg: 0.0158 (0.0257) time: 0.4098 data: 0.0083 max
mem: 6183
Epoch: [1] Total time: 0:00:25 (0.4228 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24 model_time: 0.1586 (0.1586)
evaluator_time: 0.0226 (0.0226) time: 0.4891 data: 0.3065 max mem:
6183
Test: [49/50] eta: 0:00:00 model_time: 0.0857 (0.0932)
evaluator_time: 0.0095 (0.0121) time: 0.1026 data: 0.0037 max mem:
6183
Test: Total time: 0:00:06 (0.1200 s / it)
Averaged stats: model_time: 0.0857 (0.0932) evaluator_time: 0.0095
(0.0121)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
```

```

DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.187
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.545
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.025
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.001
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.199
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.110
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.376
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.390
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.014
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.413
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.195
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.625
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.017
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.218
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.120
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.308
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.318
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.043

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.335
Epoch: [2] [ 0/60] eta: 0:01:00 lr: 0.005000 loss: 1.1471 (1.1471)
loss_classifier: 0.1634 (0.1634) loss_box_reg: 0.1893 (0.1893)
loss_mask: 0.6610 (0.6610) loss_objectness: 0.0976 (0.0976)
loss_rpn_box_reg: 0.0358 (0.0358) time: 1.0152 data: 0.3117 max
mem: 6183
Epoch: [2] [10/60] eta: 0:00:24 lr: 0.005000 loss: 0.7911 (0.8399)
loss_classifier: 0.1262 (0.1248) loss_box_reg: 0.1725 (0.1754)
loss_mask: 0.4143 (0.4348) loss_objectness: 0.0736 (0.0780)
loss_rpn_box_reg: 0.0235 (0.0269) time: 0.4916 data: 0.0387 max
mem: 6183
Epoch: [2] [20/60] eta: 0:00:19 lr: 0.005000 loss: 0.7463 (0.7800)
loss_classifier: 0.0967 (0.1107) loss_box_reg: 0.1487 (0.1608)
loss_mask: 0.3866 (0.4070) loss_objectness: 0.0597 (0.0760)
loss_rpn_box_reg: 0.0211 (0.0256) time: 0.4530 data: 0.0129 max
mem: 6183
Epoch: [2] [30/60] eta: 0:00:13 lr: 0.005000 loss: 0.7201 (0.7765)
loss_classifier: 0.0961 (0.1138) loss_box_reg: 0.1487 (0.1704)
loss_mask: 0.3622 (0.3954) loss_objectness: 0.0534 (0.0705)
loss_rpn_box_reg: 0.0237 (0.0265) time: 0.4463 data: 0.0119 max
mem: 6183
Epoch: [2] [40/60] eta: 0:00:09 lr: 0.005000 loss: 0.7201 (0.7566)
loss_classifier: 0.0987 (0.1097) loss_box_reg: 0.1438 (0.1641)
loss_mask: 0.3633 (0.3919) loss_objectness: 0.0515 (0.0656)
loss_rpn_box_reg: 0.0239 (0.0253) time: 0.4342 data: 0.0099 max
mem: 6183
Epoch: [2] [50/60] eta: 0:00:04 lr: 0.005000 loss: 0.7464 (0.7527)
loss_classifier: 0.0987 (0.1097) loss_box_reg: 0.1450 (0.1646)
loss_mask: 0.3633 (0.3901) loss_objectness: 0.0465 (0.0625)
loss_rpn_box_reg: 0.0239 (0.0259) time: 0.4306 data: 0.0099 max
mem: 6183
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.7483 (0.7566)
loss_classifier: 0.1072 (0.1100) loss_box_reg: 0.1488 (0.1696)
loss_mask: 0.3710 (0.3897) loss_objectness: 0.0465 (0.0600)
loss_rpn_box_reg: 0.0305 (0.0274) time: 0.4229 data: 0.0085 max
mem: 6183
Epoch: [2] Total time: 0:00:26 (0.4492 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:36 model_time: 0.2338 (0.2338)
evaluator_time: 0.0356 (0.0356) time: 0.7262 data: 0.4553 max mem:
6183
Test: [49/50] eta: 0:00:00 model_time: 0.1016 (0.1135)
evaluator_time: 0.0135 (0.0187) time: 0.1327 data: 0.0037 max mem:
6183
Test: Total time: 0:00:07 (0.1503 s / it)
Averaged stats: model_time: 0.1016 (0.1135) evaluator_time: 0.0135
(0.0187)

```

Accumulating evaluation results...

DONE (t=0.02s).

Accumulating evaluation results...

DONE (t=0.03s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all |  
maxDets=100 ] = 0.172

Average Precision (AP) @[ IoU=0.50 | area= all |  
maxDets=100 ] = 0.541

Average Precision (AP) @[ IoU=0.75 | area= all |  
maxDets=100 ] = 0.033

Average Precision (AP) @[ IoU=0.50:0.95 | area= small |  
maxDets=100 ] = -1.000

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |  
maxDets=100 ] = 0.004

Average Precision (AP) @[ IoU=0.50:0.95 | area= large |  
maxDets=100 ] = 0.186

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=  
1 ] = 0.087

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=  
10 ] = 0.413

Average Recall (AR) @[ IoU=0.50:0.95 | area= all |  
maxDets=100 ] = 0.426

Average Recall (AR) @[ IoU=0.50:0.95 | area= small |  
maxDets=100 ] = -1.000

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |  
maxDets=100 ] = 0.029

Average Recall (AR) @[ IoU=0.50:0.95 | area= large |  
maxDets=100 ] = 0.451

IoU metric: segm

Average Precision (AP) @[ IoU=0.50:0.95 | area= all |  
maxDets=100 ] = 0.229

Average Precision (AP) @[ IoU=0.50 | area= all |  
maxDets=100 ] = 0.662

Average Precision (AP) @[ IoU=0.75 | area= all |  
maxDets=100 ] = 0.061

Average Precision (AP) @[ IoU=0.50:0.95 | area= small |  
maxDets=100 ] = -1.000

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |  
maxDets=100 ] = 0.005

Average Precision (AP) @[ IoU=0.50:0.95 | area= large |  
maxDets=100 ] = 0.253

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=  
1 ] = 0.142

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=  
10 ] = 0.366

Average Recall (AR) @[ IoU=0.50:0.95 | area= all |  
maxDets=100 ] = 0.378

Average Recall (AR) @[ IoU=0.50:0.95 | area= small |



```
maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.114
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.395
Epoch: [3] [ 0/60] eta: 0:00:52 lr: 0.000500 loss: 0.8332 (0.8332)
loss_classifier: 0.1469 (0.1469) loss_box_reg: 0.2316 (0.2316)
loss_mask: 0.3584 (0.3584) loss_objectness: 0.0625 (0.0625)
loss_rpn_box_reg: 0.0337 (0.0337) time: 0.8698 data: 0.3433 max
mem: 6183
Epoch: [3] [10/60] eta: 0:00:23 lr: 0.000500 loss: 0.6265 (0.6456)
loss_classifier: 0.0793 (0.0893) loss_box_reg: 0.1117 (0.1371)
loss_mask: 0.3481 (0.3566) loss_objectness: 0.0410 (0.0410)
loss_rpn_box_reg: 0.0198 (0.0216) time: 0.4714 data: 0.0408 max
mem: 6183
Epoch: [3] [20/60] eta: 0:00:18 lr: 0.000500 loss: 0.7121 (0.7147)
loss_classifier: 0.0970 (0.1055) loss_box_reg: 0.1546 (0.1658)
loss_mask: 0.3503 (0.3704) loss_objectness: 0.0410 (0.0477)
loss_rpn_box_reg: 0.0236 (0.0254) time: 0.4460 data: 0.0113 max
mem: 6183
Epoch: [3] [30/60] eta: 0:00:13 lr: 0.000500 loss: 0.7439 (0.7137)
loss_classifier: 0.1094 (0.1074) loss_box_reg: 0.1584 (0.1659)
loss_mask: 0.3725 (0.3671) loss_objectness: 0.0466 (0.0475)
loss_rpn_box_reg: 0.0239 (0.0258) time: 0.4461 data: 0.0108 max
mem: 6183
Epoch: [3] [40/60] eta: 0:00:08 lr: 0.000500 loss: 0.6234 (0.6886)
loss_classifier: 0.0945 (0.1037) loss_box_reg: 0.1393 (0.1584)
loss_mask: 0.3224 (0.3561) loss_objectness: 0.0407 (0.0465)
loss_rpn_box_reg: 0.0193 (0.0238) time: 0.4300 data: 0.0094 max
mem: 6183
Epoch: [3] [50/60] eta: 0:00:04 lr: 0.000500 loss: 0.5540 (0.6643)
loss_classifier: 0.0719 (0.0975) loss_box_reg: 0.0995 (0.1491)
loss_mask: 0.3218 (0.3502) loss_objectness: 0.0392 (0.0457)
loss_rpn_box_reg: 0.0130 (0.0218) time: 0.4317 data: 0.0108 max
mem: 6183
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.5885 (0.6620)
loss_classifier: 0.0780 (0.0969) loss_box_reg: 0.1203 (0.1482)
loss_mask: 0.3260 (0.3488) loss_objectness: 0.0382 (0.0466)
loss_rpn_box_reg: 0.0165 (0.0215) time: 0.4255 data: 0.0102 max
mem: 6183
Epoch: [3] Total time: 0:00:26 (0.4436 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24 model_time: 0.1375 (0.1375)
evaluator_time: 0.0135 (0.0135) time: 0.4847 data: 0.3322 max mem:
6183
Test: [49/50] eta: 0:00:00 model_time: 0.1000 (0.1176)
evaluator_time: 0.0107 (0.0175) time: 0.1229 data: 0.0038 max mem:
6183
```

```
Test: Total time: 0:00:07 (0.1531 s / it)
Averaged stats: model_time: 0.1000 (0.1176) evaluator_time: 0.0107
(0.0175)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.321
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.691
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.228
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.002
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.343
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.157
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.495
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.503
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.014
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.534
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.266
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.701
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.084
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.007
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.298
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.149
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.367
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.370
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.086
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.388
Epoch: [4] [ 0/60] eta: 0:01:00 lr: 0.000500 loss: 0.8458 (0.8458)
loss_classifier: 0.1525 (0.1525) loss_box_reg: 0.2353 (0.2353)
loss_mask: 0.3796 (0.3796) loss_objectness: 0.0449 (0.0449)
loss_rpn_box_reg: 0.0335 (0.0335) time: 1.0156 data: 0.4515 max
mem: 6183
Epoch: [4] [10/60] eta: 0:00:23 lr: 0.000500 loss: 0.6267 (0.6354)
loss_classifier: 0.0832 (0.0913) loss_box_reg: 0.1378 (0.1425)
loss_mask: 0.3424 (0.3462) loss_objectness: 0.0352 (0.0376)
loss_rpn_box_reg: 0.0161 (0.0179) time: 0.4774 data: 0.0477 max
mem: 6183
Epoch: [4] [20/60] eta: 0:00:18 lr: 0.000500 loss: 0.6267 (0.6643)
loss_classifier: 0.0832 (0.0973) loss_box_reg: 0.1378 (0.1537)
loss_mask: 0.3354 (0.3528) loss_objectness: 0.0357 (0.0426)
loss_rpn_box_reg: 0.0145 (0.0179) time: 0.4387 data: 0.0098 max
mem: 6183
Epoch: [4] [30/60] eta: 0:00:13 lr: 0.000500 loss: 0.6212 (0.6537)
loss_classifier: 0.0890 (0.0946) loss_box_reg: 0.1292 (0.1488)
loss_mask: 0.3219 (0.3465) loss_objectness: 0.0441 (0.0441)
loss_rpn_box_reg: 0.0165 (0.0196) time: 0.4395 data: 0.0110 max
mem: 6183
Epoch: [4] [40/60] eta: 0:00:08 lr: 0.000500 loss: 0.5359 (0.6376)
loss_classifier: 0.0712 (0.0911) loss_box_reg: 0.1084 (0.1442)
loss_mask: 0.3119 (0.3369) loss_objectness: 0.0391 (0.0457)
loss_rpn_box_reg: 0.0183 (0.0197) time: 0.4222 data: 0.0092 max
mem: 6183
Epoch: [4] [50/60] eta: 0:00:04 lr: 0.000500 loss: 0.5708 (0.6422)
loss_classifier: 0.0751 (0.0916) loss_box_reg: 0.1189 (0.1453)
loss_mask: 0.3252 (0.3399) loss_objectness: 0.0331 (0.0447)
loss_rpn_box_reg: 0.0198 (0.0207) time: 0.4346 data: 0.0093 max
mem: 6183
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.6419 (0.6426)
loss_classifier: 0.0798 (0.0911) loss_box_reg: 0.1289 (0.1465)
loss_mask: 0.3501 (0.3412) loss_objectness: 0.0314 (0.0432)
loss_rpn_box_reg: 0.0219 (0.0207) time: 0.4379 data: 0.0092 max
mem: 6183
Epoch: [4] Total time: 0:00:26 (0.4462 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24 model_time: 0.1605 (0.1605)
evaluator_time: 0.0165 (0.0165) time: 0.4862 data: 0.3077 max mem:
6183

```

```
Test: [49/50] eta: 0:00:00 model_time: 0.1222 (0.1189)
evaluator_time: 0.0208 (0.0195) time: 0.1653 data: 0.0063 max mem:
6183
Test: Total time: 0:00:07 (0.1563 s / it)
Averaged stats: model_time: 0.1222 (0.1189) evaluator_time: 0.0208
(0.0195)
Accumulating evaluation results...
DONE (t=0.03s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.308
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.698
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.206
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.007
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.327
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.164
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.502
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.505
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.029
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.535
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.262
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.667
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.080
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.003
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.285
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.163
```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.370
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.376
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.071
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.395
Epoch: [5] [ 0/60] eta: 0:01:03  lr: 0.000500  loss: 0.7245 (0.7245)
loss_classifier: 0.1085 (0.1085)  loss_box_reg: 0.1718 (0.1718)
loss_mask: 0.3913 (0.3913)  loss_objectness: 0.0277 (0.0277)
loss_rpn_box_reg: 0.0250 (0.0250)  time: 1.0635  data: 0.4115  max
mem: 6183
Epoch: [5] [10/60] eta: 0:00:24  lr: 0.000500  loss: 0.6815 (0.7052)
loss_classifier: 0.1052 (0.1035)  loss_box_reg: 0.1591 (0.1632)
loss_mask: 0.3807 (0.3729)  loss_objectness: 0.0422 (0.0459)
loss_rpn_box_reg: 0.0184 (0.0196)  time: 0.4903  data: 0.0447  max
mem: 6183
Epoch: [5] [20/60] eta: 0:00:18  lr: 0.000500  loss: 0.6325 (0.6620)
loss_classifier: 0.0885 (0.0937)  loss_box_reg: 0.1364 (0.1454)
loss_mask: 0.3398 (0.3554)  loss_objectness: 0.0422 (0.0444)
loss_rpn_box_reg: 0.0203 (0.0231)  time: 0.4343  data: 0.0090  max
mem: 6183
Epoch: [5] [30/60] eta: 0:00:13  lr: 0.000500  loss: 0.5582 (0.6269)
loss_classifier: 0.0639 (0.0864)  loss_box_reg: 0.1032 (0.1318)
loss_mask: 0.3040 (0.3465)  loss_objectness: 0.0346 (0.0419)
loss_rpn_box_reg: 0.0174 (0.0204)  time: 0.4385  data: 0.0100  max
mem: 6183
Epoch: [5] [40/60] eta: 0:00:08  lr: 0.000500  loss: 0.5365 (0.6248)
loss_classifier: 0.0671 (0.0859)  loss_box_reg: 0.1032 (0.1328)
loss_mask: 0.3096 (0.3436)  loss_objectness: 0.0346 (0.0422)
loss_rpn_box_reg: 0.0140 (0.0203)  time: 0.4347  data: 0.0098  max
mem: 6183
Epoch: [5] [50/60] eta: 0:00:04  lr: 0.000500  loss: 0.6168 (0.6271)
loss_classifier: 0.0794 (0.0864)  loss_box_reg: 0.1335 (0.1345)
loss_mask: 0.3255 (0.3438)  loss_objectness: 0.0386 (0.0422)
loss_rpn_box_reg: 0.0183 (0.0202)  time: 0.4273  data: 0.0098  max
mem: 6183
Epoch: [5] [59/60] eta: 0:00:00  lr: 0.000500  loss: 0.6168 (0.6215)
loss_classifier: 0.0864 (0.0869)  loss_box_reg: 0.1335 (0.1357)
loss_mask: 0.3076 (0.3368)  loss_objectness: 0.0361 (0.0421)
loss_rpn_box_reg: 0.0190 (0.0200)  time: 0.4290  data: 0.0089  max
mem: 6183
Epoch: [5] Total time: 0:00:26 (0.4488 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24  model_time: 0.1630 (0.1630)

```

```
evaluator_time: 0.0140 (0.0140)  time: 0.4981  data: 0.3196  max mem: 6183
Test: [49/50]  eta: 0:00:00  model_time: 0.1020 (0.1077)
evaluator_time: 0.0105 (0.0132)  time: 0.1234  data: 0.0040  max mem: 6183
Test: Total time: 0:00:06 (0.1360 s / it)
Averaged stats: model_time: 0.1020 (0.1077)  evaluator_time: 0.0105 (0.0132)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.313
Average Precision  (AP) @[ IoU=0.50      | area=  all | maxDets=100 ] = 0.719
Average Precision  (AP) @[ IoU=0.75      | area=  all | maxDets=100 ] = 0.234
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.006
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.335
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets= 1 ] = 0.161
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=10 ] = 0.490
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.497
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.029
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.527
IoU metric: segm
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.278
Average Precision  (AP) @[ IoU=0.50      | area=  all | maxDets=100 ] = 0.718
Average Precision  (AP) @[ IoU=0.75      | area=  all | maxDets=100 ] = 0.123
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.008
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
```

```
maxDets=100 ] = 0.309
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.161
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.392
Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.396
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.114
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.413
Epoch: [6] [ 0/60] eta: 0:00:57 lr: 0.000050 loss: 0.7117 (0.7117)
loss_classifier: 0.1020 (0.1020) loss_box_reg: 0.2017 (0.2017)
loss_mask: 0.3243 (0.3243) loss_objectness: 0.0519 (0.0519)
loss_rpn_box_reg: 0.0317 (0.0317) time: 0.9569 data: 0.3761 max
mem: 6183
Epoch: [6] [10/60] eta: 0:00:25 lr: 0.000050 loss: 0.6789 (0.6812)
loss_classifier: 0.0925 (0.0971) loss_box_reg: 0.1653 (0.1522)
loss_mask: 0.3516 (0.3547) loss_objectness: 0.0468 (0.0554)
loss_rpn_box_reg: 0.0201 (0.0217) time: 0.5068 data: 0.0427 max
mem: 6183
Epoch: [6] [20/60] eta: 0:00:18 lr: 0.000050 loss: 0.6227 (0.6372)
loss_classifier: 0.0839 (0.0902) loss_box_reg: 0.1162 (0.1417)
loss_mask: 0.3342 (0.3348) loss_objectness: 0.0404 (0.0504)
loss_rpn_box_reg: 0.0188 (0.0201) time: 0.4423 data: 0.0091 max
mem: 6183
Epoch: [6] [30/60] eta: 0:00:13 lr: 0.000050 loss: 0.5788 (0.6376)
loss_classifier: 0.0795 (0.0867) loss_box_reg: 0.1121 (0.1373)
loss_mask: 0.3030 (0.3438) loss_objectness: 0.0400 (0.0489)
loss_rpn_box_reg: 0.0196 (0.0210) time: 0.4211 data: 0.0091 max
mem: 6183
Epoch: [6] [40/60] eta: 0:00:09 lr: 0.000050 loss: 0.5370 (0.6097)
loss_classifier: 0.0746 (0.0827) loss_box_reg: 0.1006 (0.1308)
loss_mask: 0.3002 (0.3321) loss_objectness: 0.0322 (0.0441)
loss_rpn_box_reg: 0.0200 (0.0200) time: 0.4399 data: 0.0099 max
mem: 6183
Epoch: [6] [50/60] eta: 0:00:04 lr: 0.000050 loss: 0.5223 (0.6096)
loss_classifier: 0.0680 (0.0837) loss_box_reg: 0.1117 (0.1326)
loss_mask: 0.2959 (0.3314) loss_objectness: 0.0298 (0.0428)
loss_rpn_box_reg: 0.0132 (0.0191) time: 0.4486 data: 0.0102 max
mem: 6183
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.5875 (0.6110)
loss_classifier: 0.0843 (0.0846) loss_box_reg: 0.1176 (0.1345)
loss_mask: 0.3157 (0.3312) loss_objectness: 0.0330 (0.0417)
loss_rpn_box_reg: 0.0158 (0.0189) time: 0.4291 data: 0.0089 max
mem: 6183
Epoch: [6] Total time: 0:00:26 (0.4486 s / it)
```

```
creating index...
index created!
Test: [ 0/50] eta: 0:00:37 model_time: 0.2324 (0.2324)
evaluator_time: 0.0248 (0.0248) time: 0.7444 data: 0.4851 max mem:
6183
Test: [49/50] eta: 0:00:00 model_time: 0.1017 (0.1117)
evaluator_time: 0.0113 (0.0136) time: 0.1266 data: 0.0040 max mem:
6183
Test: Total time: 0:00:07 (0.1438 s / it)
Averaged stats: model_time: 0.1017 (0.1117) evaluator_time: 0.0113
(0.0136)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.03s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.302
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.685
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.205
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.000
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.324
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.173
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.483
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.492
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.523
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.275
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.713
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.104
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
```



```
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.006
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.307
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.169
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.385
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.388
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.100
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.406
Epoch: [7] [ 0/60] eta: 0:00:57 lr: 0.000050 loss: 0.5761 (0.5761)
loss_classifier: 0.0837 (0.0837) loss_box_reg: 0.1236 (0.1236)
loss_mask: 0.3147 (0.3147) loss_objectness: 0.0353 (0.0353)
loss_rpn_box_reg: 0.0188 (0.0188) time: 0.9569 data: 0.4152 max
mem: 6183
Epoch: [7] [10/60] eta: 0:00:24 lr: 0.000050 loss: 0.5761 (0.6286)
loss_classifier: 0.0820 (0.0869) loss_box_reg: 0.1268 (0.1465)
loss_mask: 0.3147 (0.3304) loss_objectness: 0.0359 (0.0445)
loss_rpn_box_reg: 0.0174 (0.0203) time: 0.4945 data: 0.0456 max
mem: 6183
Epoch: [7] [20/60] eta: 0:00:18 lr: 0.000050 loss: 0.5666 (0.6401)
loss_classifier: 0.0784 (0.0890) loss_box_reg: 0.1336 (0.1466)
loss_mask: 0.3140 (0.3366) loss_objectness: 0.0397 (0.0465)
loss_rpn_box_reg: 0.0226 (0.0215) time: 0.4492 data: 0.0097 max
mem: 6183
Epoch: [7] [30/60] eta: 0:00:13 lr: 0.000050 loss: 0.5603 (0.6153)
loss_classifier: 0.0743 (0.0847) loss_box_reg: 0.1095 (0.1369)
loss_mask: 0.3106 (0.3308) loss_objectness: 0.0374 (0.0432)
loss_rpn_box_reg: 0.0208 (0.0197) time: 0.4340 data: 0.0095 max
mem: 6183
Epoch: [7] [40/60] eta: 0:00:09 lr: 0.000050 loss: 0.5922 (0.6248)
loss_classifier: 0.0779 (0.0881) loss_box_reg: 0.1081 (0.1423)
loss_mask: 0.3125 (0.3308) loss_objectness: 0.0364 (0.0438)
loss_rpn_box_reg: 0.0184 (0.0198) time: 0.4351 data: 0.0083 max
mem: 6183
Epoch: [7] [50/60] eta: 0:00:04 lr: 0.000050 loss: 0.6158 (0.6204)
loss_classifier: 0.0936 (0.0866) loss_box_reg: 0.1060 (0.1399)
loss_mask: 0.3191 (0.3313) loss_objectness: 0.0408 (0.0430)
loss_rpn_box_reg: 0.0184 (0.0197) time: 0.4539 data: 0.0096 max
mem: 6183
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.5649 (0.6236)
loss_classifier: 0.0800 (0.0864) loss_box_reg: 0.1205 (0.1402)
loss_mask: 0.3216 (0.3355) loss_objectness: 0.0391 (0.0422)
```

```
loss_rpn_box_reg: 0.0162 (0.0195)  time: 0.4381  data: 0.0091  max
mem: 6183
Epoch: [7] Total time: 0:00:27 (0.4523 s / it)
creating index...
index created!
Test: [ 0/50]  eta: 0:00:24  model_time: 0.1626 (0.1626)
evaluator_time: 0.0159 (0.0159)  time: 0.4915  data: 0.3013  max mem:
6183
Test: [49/50]  eta: 0:00:00  model_time: 0.1015 (0.1229)
evaluator_time: 0.0091 (0.0170)  time: 0.1260  data: 0.0040  max mem:
6183
Test: Total time: 0:00:07 (0.1579 s / it)
Averaged stats: model_time: 0.1015 (0.1229)  evaluator_time: 0.0091
(0.0170)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.307
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.710
Average Precision  (AP) @[ IoU=0.75      | area=  all |
maxDets=100 ] = 0.174
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.007
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.327
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
1 ] = 0.178
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=
10 ] = 0.482
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.494
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.043
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.522
IoU metric: segm
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all |
maxDets=100 ] = 0.279
Average Precision  (AP) @[ IoU=0.50      | area=  all |
maxDets=100 ] = 0.691
Average Precision  (AP) @[ IoU=0.75      | area=  all |
```

```

maxDets=100 ] = 0.111
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.007
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.312
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.171
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.388
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.391
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.129
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.407
Epoch: [8] [ 0/60] eta: 0:00:55 lr: 0.000050 loss: 0.6837 (0.6837)
loss_classifier: 0.1019 (0.1019) loss_box_reg: 0.1832 (0.1832)
loss_mask: 0.3461 (0.3461) loss_objectness: 0.0296 (0.0296)
loss_rpn_box_reg: 0.0229 (0.0229) time: 0.9262 data: 0.3697 max
mem: 6183
Epoch: [8] [10/60] eta: 0:00:23 lr: 0.000050 loss: 0.5116 (0.5398)
loss_classifier: 0.0680 (0.0746) loss_box_reg: 0.1039 (0.1144)
loss_mask: 0.2978 (0.2984) loss_objectness: 0.0340 (0.0349)
loss_rpn_box_reg: 0.0143 (0.0175) time: 0.4659 data: 0.0402 max
mem: 6183
Epoch: [8] [20/60] eta: 0:00:18 lr: 0.000050 loss: 0.4972 (0.5452)
loss_classifier: 0.0680 (0.0756) loss_box_reg: 0.0901 (0.1090)
loss_mask: 0.2931 (0.3038) loss_objectness: 0.0340 (0.0404)
loss_rpn_box_reg: 0.0143 (0.0164) time: 0.4350 data: 0.0083 max
mem: 6183
Epoch: [8] [30/60] eta: 0:00:13 lr: 0.000050 loss: 0.5934 (0.5805)
loss_classifier: 0.0765 (0.0822) loss_box_reg: 0.1062 (0.1195)
loss_mask: 0.3194 (0.3211) loss_objectness: 0.0365 (0.0406)
loss_rpn_box_reg: 0.0148 (0.0171) time: 0.4453 data: 0.0089 max
mem: 6183
Epoch: [8] [40/60] eta: 0:00:09 lr: 0.000050 loss: 0.6306 (0.6022)
loss_classifier: 0.0970 (0.0869) loss_box_reg: 0.1469 (0.1339)
loss_mask: 0.3323 (0.3207) loss_objectness: 0.0382 (0.0420)
loss_rpn_box_reg: 0.0218 (0.0188) time: 0.4470 data: 0.0085 max
mem: 6244
Epoch: [8] [50/60] eta: 0:00:04 lr: 0.000050 loss: 0.6306 (0.6087)
loss_classifier: 0.0890 (0.0869) loss_box_reg: 0.1527 (0.1359)
loss_mask: 0.3323 (0.3256) loss_objectness: 0.0343 (0.0415)
loss_rpn_box_reg: 0.0217 (0.0189) time: 0.4622 data: 0.0119 max
mem: 6244

```

```
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.6214 (0.6103)
loss_classifier: 0.0828 (0.0869) loss_box_reg: 0.1358 (0.1360)
loss_mask: 0.3344 (0.3274) loss_objectness: 0.0327 (0.0411)
loss_rpn_box_reg: 0.0187 (0.0189) time: 0.4503 data: 0.0115 max
mem: 6244
Epoch: [8] Total time: 0:00:27 (0.4548 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24 model_time: 0.2058 (0.2058)
evaluator_time: 0.0193 (0.0193) time: 0.4965 data: 0.2696 max mem:
6244
Test: [49/50] eta: 0:00:00 model_time: 0.1229 (0.1261)
evaluator_time: 0.0158 (0.0189) time: 0.1597 data: 0.0079 max mem:
6244
Test: Total time: 0:00:08 (0.1621 s / it)
Averaged stats: model_time: 0.1229 (0.1261) evaluator_time: 0.0158
(0.0189)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.337
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.702
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.241
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.002
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.361
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.187
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.518
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.528
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.014
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.560
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.274
```

```

Average Precision (AP) @[ IoU=0.50      | area=   all |
maxDets=100 ] = 0.691
Average Precision (AP) @[ IoU=0.75      | area=   all |
maxDets=100 ] = 0.111
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.007
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.303
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
1 ] = 0.163
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=
10 ] = 0.382
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all |
maxDets=100 ] = 0.383
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.143
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.398
Epoch: [9] [ 0/60] eta: 0:00:55 lr: 0.000005 loss: 0.5576 (0.5576)
loss_classifier: 0.0978 (0.0978) loss_box_reg: 0.1063 (0.1063)
loss_mask: 0.2981 (0.2981) loss_objectness: 0.0373 (0.0373)
loss_rpn_box_reg: 0.0181 (0.0181) time: 0.9230 data: 0.3424 max
mem: 6244
Epoch: [9] [10/60] eta: 0:00:23 lr: 0.000005 loss: 0.5576 (0.6111)
loss_classifier: 0.0718 (0.0837) loss_box_reg: 0.1093 (0.1259)
loss_mask: 0.3419 (0.3431) loss_objectness: 0.0373 (0.0416)
loss_rpn_box_reg: 0.0177 (0.0169) time: 0.4736 data: 0.0390 max
mem: 6244
Epoch: [9] [20/60] eta: 0:00:18 lr: 0.000005 loss: 0.5708 (0.6228)
loss_classifier: 0.0752 (0.0859) loss_box_reg: 0.1188 (0.1423)
loss_mask: 0.3121 (0.3319) loss_objectness: 0.0370 (0.0423)
loss_rpn_box_reg: 0.0177 (0.0203) time: 0.4455 data: 0.0093 max
mem: 6244
Epoch: [9] [30/60] eta: 0:00:13 lr: 0.000005 loss: 0.5708 (0.6139)
loss_classifier: 0.0761 (0.0865) loss_box_reg: 0.1278 (0.1397)
loss_mask: 0.2954 (0.3253) loss_objectness: 0.0350 (0.0425)
loss_rpn_box_reg: 0.0193 (0.0198) time: 0.4619 data: 0.0098 max
mem: 6244
Epoch: [9] [40/60] eta: 0:00:09 lr: 0.000005 loss: 0.5281 (0.6200)
loss_classifier: 0.0808 (0.0855) loss_box_reg: 0.0889 (0.1376)
loss_mask: 0.3113 (0.3339) loss_objectness: 0.0355 (0.0439)
loss_rpn_box_reg: 0.0201 (0.0192) time: 0.4407 data: 0.0087 max
mem: 6244
Epoch: [9] [50/60] eta: 0:00:04 lr: 0.000005 loss: 0.6068 (0.6288)
loss_classifier: 0.0847 (0.0870) loss_box_reg: 0.1374 (0.1424)

```

```
loss_mask: 0.3301 (0.3357) loss_objectness: 0.0403 (0.0443)
loss_rpn_box_reg: 0.0197 (0.0193) time: 0.4364 data: 0.0092 max
mem: 6244
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.5627 (0.6184)
loss_classifier: 0.0730 (0.0846) loss_box_reg: 0.1200 (0.1382)
loss_mask: 0.3146 (0.3338) loss_objectness: 0.0403 (0.0433)
loss_rpn_box_reg: 0.0137 (0.0186) time: 0.4404 data: 0.0092 max
mem: 6244
Epoch: [9] Total time: 0:00:27 (0.4543 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.1727 (0.1727)
evaluator_time: 0.0176 (0.0176) time: 0.5005 data: 0.3087 max mem:
6244
Test: [49/50] eta: 0:00:00 model_time: 0.1152 (0.1169)
evaluator_time: 0.0156 (0.0170) time: 0.1473 data: 0.0044 max mem:
6244
Test: Total time: 0:00:07 (0.1522 s / it)
Averaged stats: model_time: 0.1152 (0.1169) evaluator_time: 0.0156
(0.0170)
Accumulating evaluation results...
DONE (t=0.04s).
Accumulating evaluation results...
DONE (t=0.04s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.324
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.718
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.242
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.009
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.345
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.165
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.507
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.519
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.043
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.549
```

```

IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.275
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.730
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.092
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.009
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.305
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.161
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.380
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.382
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.143
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = 0.397
That's it!

```

### Comments on training log:

In the above training log, note the last batch of the 10th epoch result.

```

Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.5627 (0.6184) loss_classifier: 0.0730
(0.0846) loss_box_reg: 0.1200 (0.1382) loss_mask: 0.3146 (0.3338) loss_objectness: 0.0403
(0.0433) loss_rpn_box_reg: 0.0137 (0.0186) time: 0.4404 data: 0.0092 max mem: 6244

```

Here, the loss value in paranthesis represents the cumulative loss over the entire epoch upto that point( here its the last batch so its for the entire epoch ) using the weights after the completion of the 10th epoch.

Similarly for loss\_classifier, loss\_mask, loss\_objectness, loss\_rpn\_box\_reg

These results can be used for comparing the model performance asked in Q5B)

### Tetsing of Backbone model on Beatles\_Abbey\_Road test image (Method 1)

```

import matplotlib.pyplot as plt
import cv2
from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

# Read an image from a specified path

```

```

image = read_image("/content/sample_data/Beatles_-_Abbey_Road.jpeg")

# Create an output image to visualize the results
output_image = image

# Obtain an evaluation transformation with 'train=False'
eval_transform = get_transform(train=False)

# Set the model in evaluation mode
model.eval()

with torch.no_grad():
    x = eval_transform(image)

    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)

    # Make predictions using the model
    predictions = model([x, ])
    pred = predictions[0]

# Normalize and convert the image to 8-bit integers (uint8)
image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]

# Filter predictions based on confidence scores (only keep scores >
0.65)
# mask refers to binary-mask ie true, false of predictions above
confidence( not meaning the mask displayed in the image)
mask = pred["scores"] > 0.65
filtered_pred = {key: value[mask] for key, value in pred.items()}

#Obtaining labels, boxes and masks for filtered predictions
filtered_labels = [f"ped: {score:.3f}" for score in
filtered_pred["scores"]]
filtered_boxes = filtered_pred["boxes"].long()
masks = (filtered_pred["masks"] > 0.7).squeeze(1)

#output image having the filtered prediction masks now
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

# Convert to NumPy array
output_image = output_image.permute(1, 2, 0).cpu().numpy().copy()

#Drawing the boxes, labels using cv2
for label, box in zip(filtered_labels, filtered_boxes):
    x_1, y_1, x_2, y_2 = [coord.item() for coord in box]


```



```
output_image = cv2.rectangle(output_image, (x_1, y_1), (x_2, y_2),  
(255, 0, 0), 2)  
output_image = cv2.putText(output_image, label, (x_1, y_1 - 10),  
cv2.FONT_HERSHEY_SIMPLEX, 0.3, (255, 0, 0), 1)
```

```
#Plotting the final image  
plt.figure(figsize=(10, 10))  
plt.imshow(output_image)
```

<matplotlib.image.AxesImage at 0x7a4c543b09a0>



```
print(pred["scores"])

tensor([0.7383, 0.5354, 0.5029, 0.4548, 0.3040, 0.2587, 0.0986,
        0.0715, 0.0707],
        device='cuda:0')
```

### Comments on testing (method 1) result:

In the above method for testing, predictions were inferred from the model. Each prediction dictionary consisted of labels, boxes, masks, confidence scores as keys.

Using confidence score threshold of 0.7, predictions in the list were filtered and the corresponding boxes, labels, masks for the the filtered predictions were outputed. In this case, of all prediction scores (as printed above) only one of them crossed the threshold.

### Tetsing of Backbone model on Beatles\_Abbey\_Road test image (Method 2)

```
import matplotlib.pyplot as plt

from torchvision.utils import draw_bounding_boxes,
draw_segmentation_masks

# Read an image from a specified path
image = read_image("/content/sample_data/Beatles_-_Abbey_Road.jpeg")

# Create an output image to visualize the results
output_image = image

# Obtain an evaluation transformation with 'train=False'
eval_transform = get_transform(train=False)

# Set the model in evaluation mode
model.eval()

with torch.no_grad():
    x = eval_transform(image)

    # convert RGBA -> RGB and move to device
    x = x[:3, ...].to(device)

    # Make predictions using the model
    predictions = model([x, ])
    pred = predictions[0]


# Normalize and convert the image to 8-bit integers (uint8)
image = (255.0 * (image - image.min()) / (image.max() -
image.min())).to(torch.uint8)
image = image[:3, ...]

#Filtering masks based on confidence
masks = (pred["masks"] > 0.7).squeeze(1)
```

```
#output image having the filtered prediction masks now
output_image = draw_segmentation_masks(output_image, masks, alpha=0.5,
colors="blue")

#Plotting the final output image
plt.figure(figsize=(12, 12))
plt.imshow(output_image.permute(1, 2, 0))

<matplotlib.image.AxesImage at 0x7a4c98c92050>
```



**Comments on testing (method 2) result:**

In the above method for testing, predictions were inferred from the model.

From each prediction, the corresponding masks whose values exceeded the set threshold were used for the output image.

Prediction dictionaries with low scores like 0.1 etc might have some values in their masks as 0.8 etc ( above the threshod). These masks are included in this method but discared in the previous method.

### **Q5.b)**

#### **Performance of two models on the training data after 10 epochs:**

*Option 1 model backbone resnet performance:*

Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1646 (0.1758) loss\_classifier: 0.0197 (0.0221) loss\_box\_reg: 0.0283 (0.0343) loss\_mask: 0.1131 (0.1162) loss\_objectness: 0.0003 (0.0008) loss\_rpn\_box\_reg: 0.0018 (0.0024) time: 0.6305 data: 0.0085 max mem: 3779

*Option 2 model backbone Mobilenet performance:*

Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.5627 (0.6184) loss\_classifier: 0.0730 (0.0846) loss\_box\_reg: 0.1200 (0.1382) loss\_mask: 0.3146 (0.3338) loss\_objectness: 0.0403 (0.0433) loss\_rpn\_box\_reg: 0.0137 (0.0186) time: 0.4404 data: 0.0092 max mem: 6244

Comparing the loss in paranthesis which represents the training loss of the entire dataset using the weights at the end of the 10th epoch, we can see option 1 model has loss of 0.1758 but option 2 has a loss of 0.6184, similarly loss\_classifier for option 1 is 0.0221 but for option 2 is 0.0846. Similarly see for loss\_mask, loss\_rpn\_box etc

**Thus option 1 model ( resnet backbone ) has lesser training loss than option 2 model ( mobilenet backbone ).**

Another way of comparing the performance is the training time.

*Option 1 model:*

Epoch: [9] Total time: 0:00:35 (0.5994 s / it)

*Option 2 model:*

Epoch: [9] Total time: 0:00:27 (0.4543 s / it)

In general, in the two model's training logs we can see that option 1 model takes more training time per epoch compared to option 2 model.

**Thus to summarize, Option 1 model ( Resnet ) has more accuracy than option 2 model ( Mobilenet ). But Option 1 Model ( Resnet ) takes more training time compared to option 2 model ( MobileNet )**

### **Q5C)**

*Testing method 1:*

In testing method 1, we saw that for option 1 model, four predictions crossed the set threshold and correspondingly four pedestrians were detected.

However, for option 2 model, only one prediction crossed the set threshold and correspondingly only one pedestrian was detected, thus missing the remaining pedestrians.

The bounding box co-ordinates and labels scores prediction in option 1 model is also better as the co-ordinates of the box are more or less bounding the pedestrian correctly that too with a high confidence score.

#### *Testing method 2:*

The masks for option 1 model is almost correctly segmenting the four pedestrians

Whereas in option 2 model, some parts of the pedestrian like their legs etc are not masked or some parts of the road is also incorrectly masked.

**To summarize, in terms of accuracy, option 1 Model (Resnet) is performing better than option 2 Model (Mobilenet) on the test image.**